



EECS 442 – Computer vision

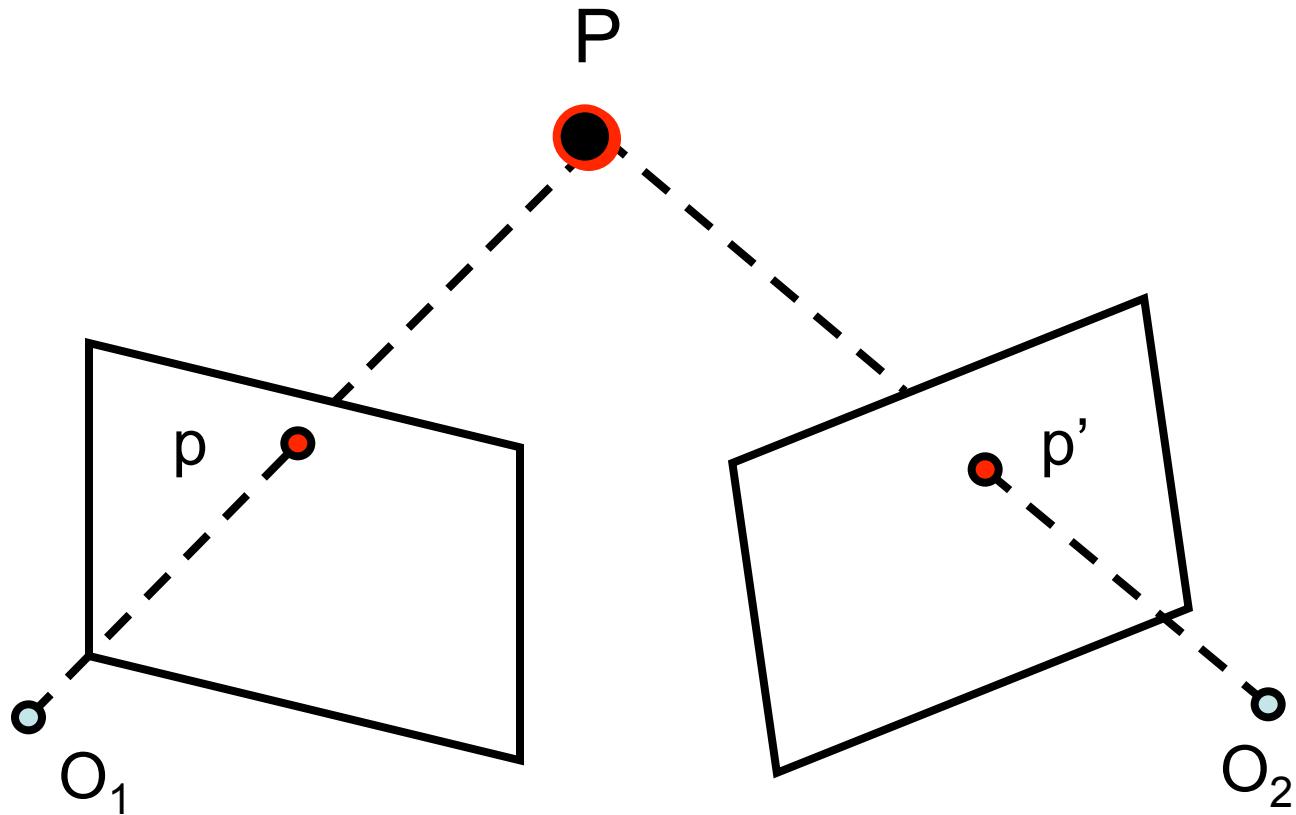
Stereo systems

- Stereo vision
- Rectification
- Correspondence problem
- Active stereo vision systems

Reading: [HZ] Chapter: 11

[FP] Chapter: 11

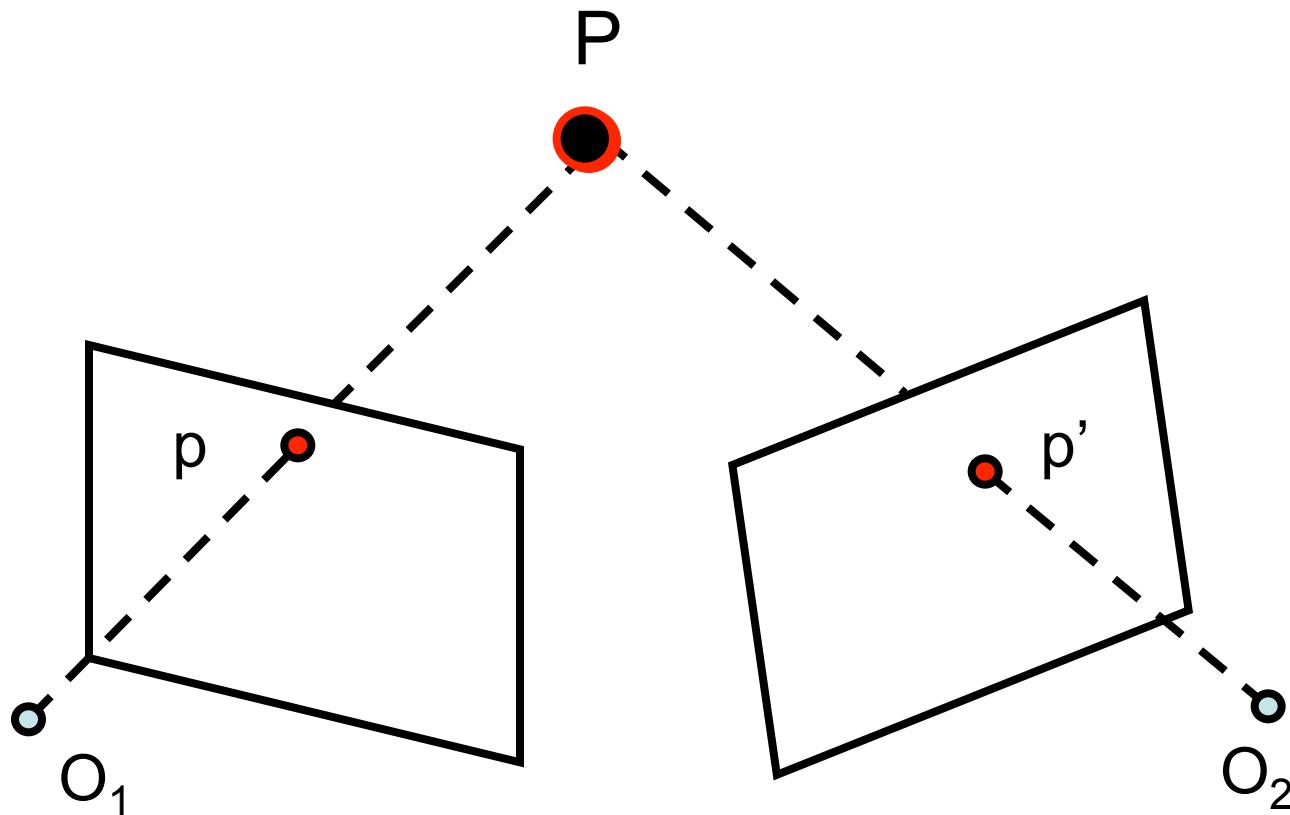
Stereo vision



Goal: estimate the position of P given the observation of P from two view points

Assumptions: known camera parameters and position (K, R, T)

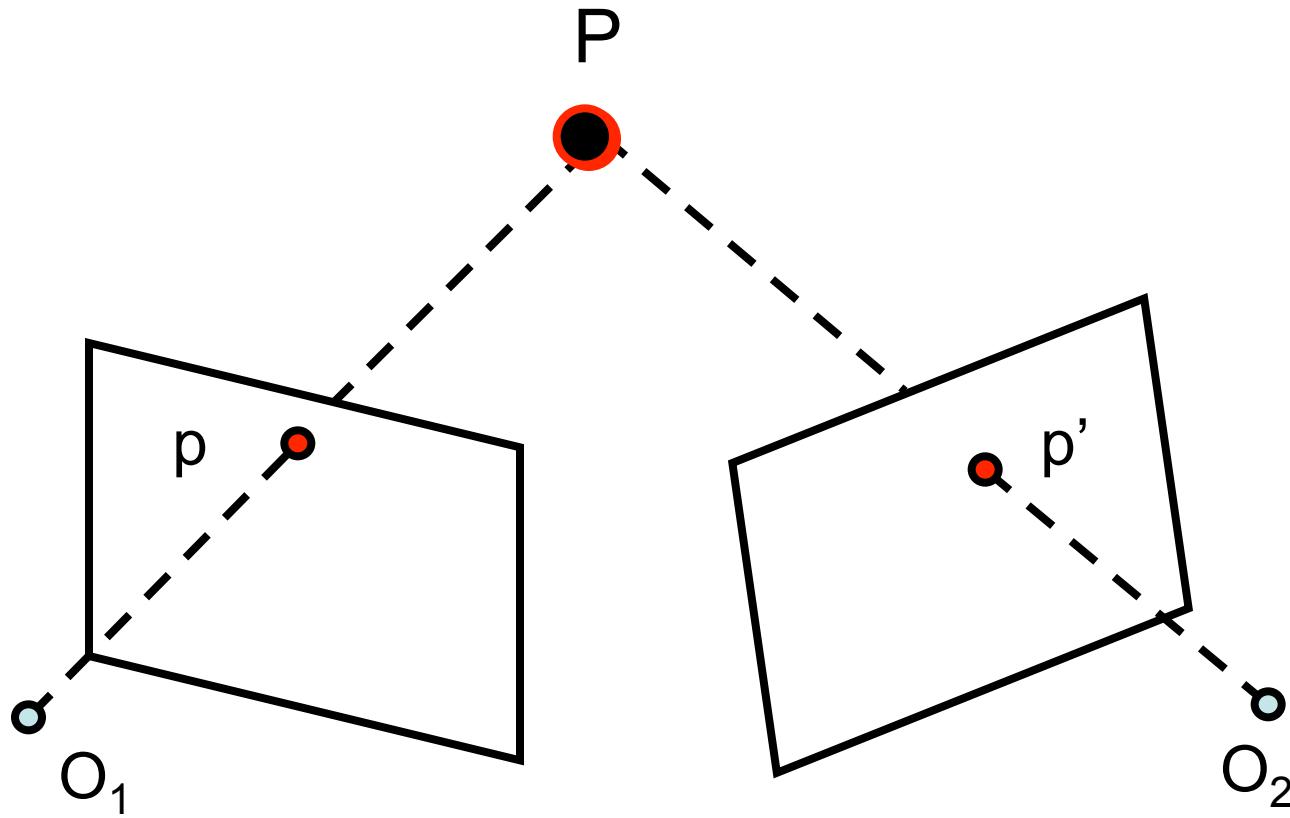
Stereo vision



Subgoals:

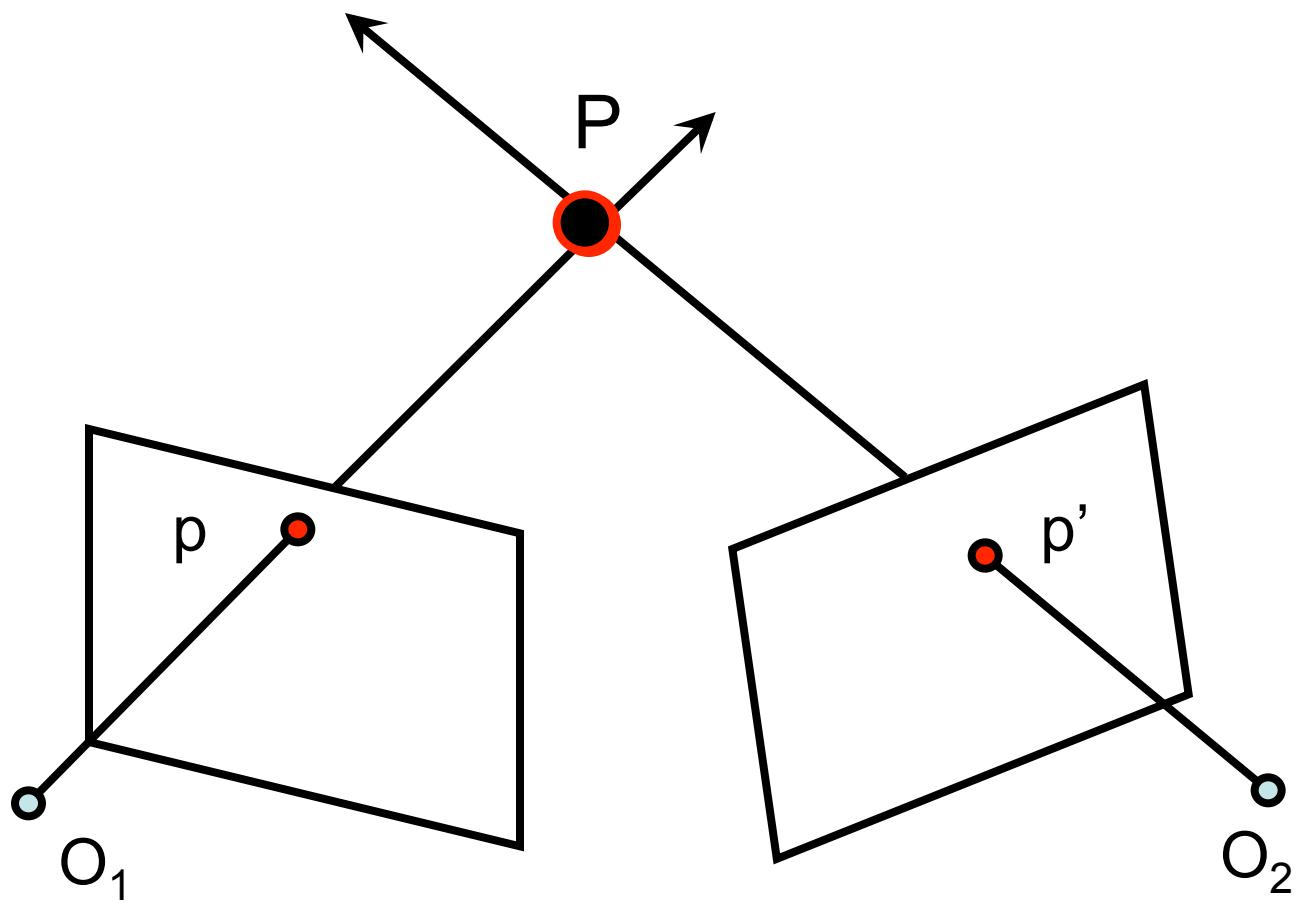
- Solve the correspondence problem
- Use corresponding observations to triangulate

Correspondence problem



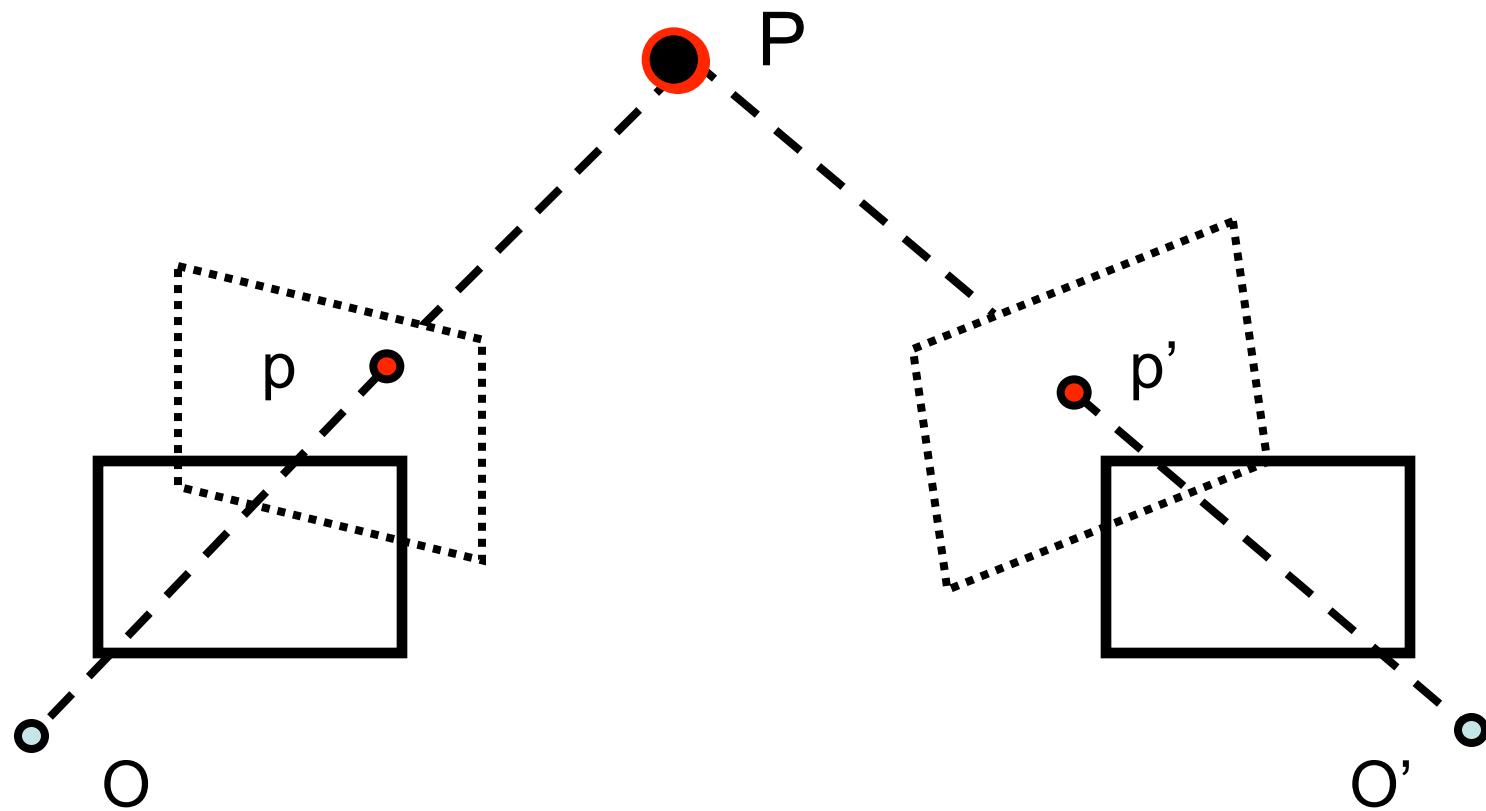
- Given a point in 3d, discover corresponding observations in left and right images

Triangulation



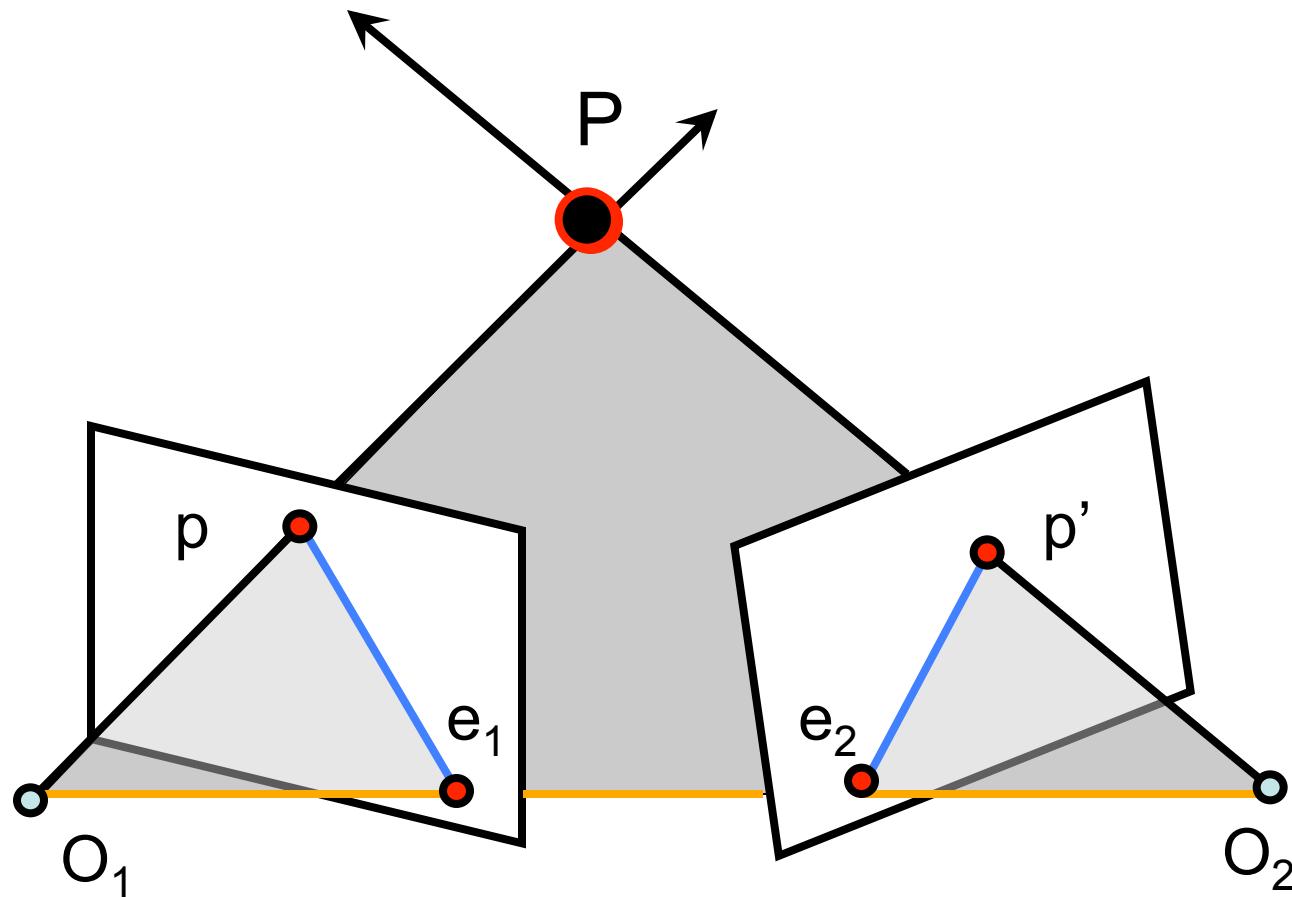
- Intersecting the two lines of sight gives rise to P

Parallel image planes



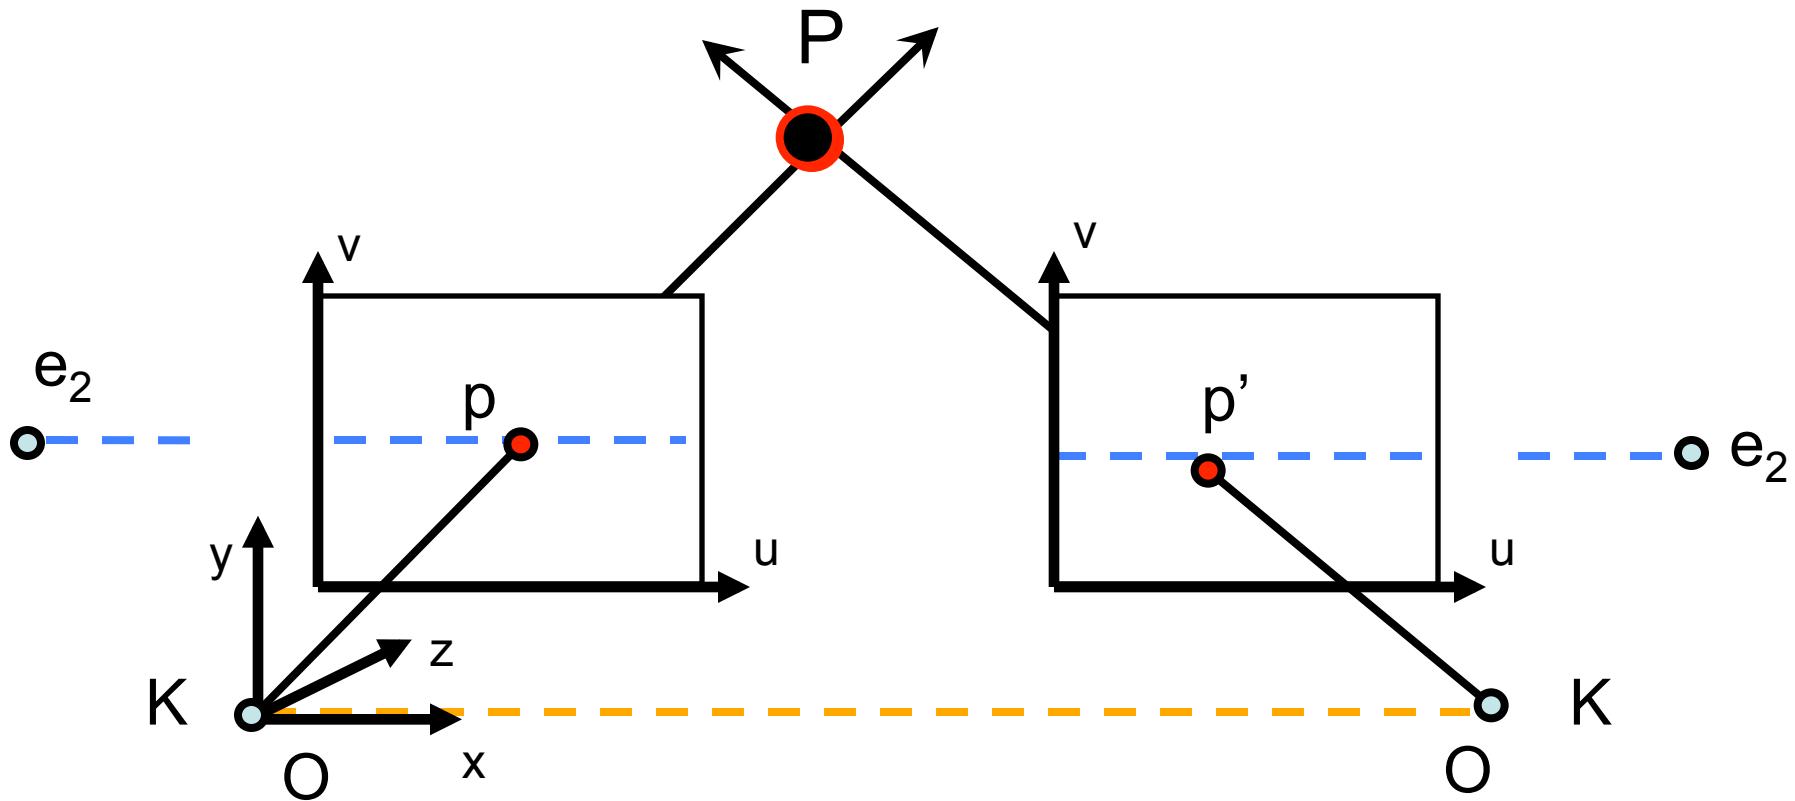
- When views are **parallel** these two steps becomes much easier!

Epipolar geometry



- Epipolar Plane
- Baseline
- Epipolar Lines
- Epipoles e_1, e_2
 - = intersections of baseline with image planes
 - = projections of the other camera center

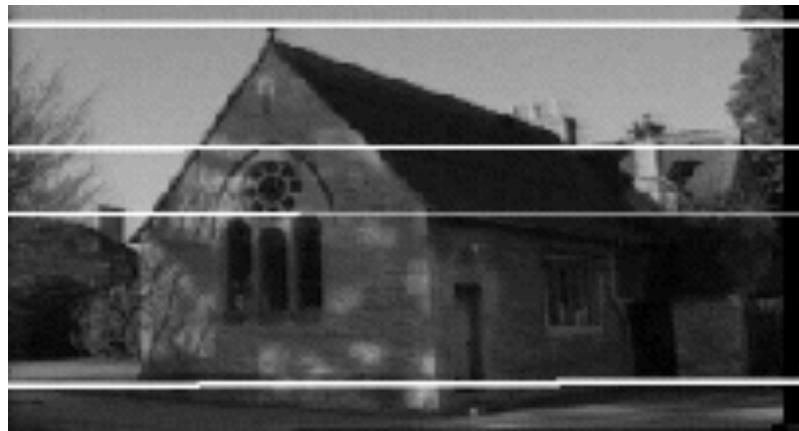
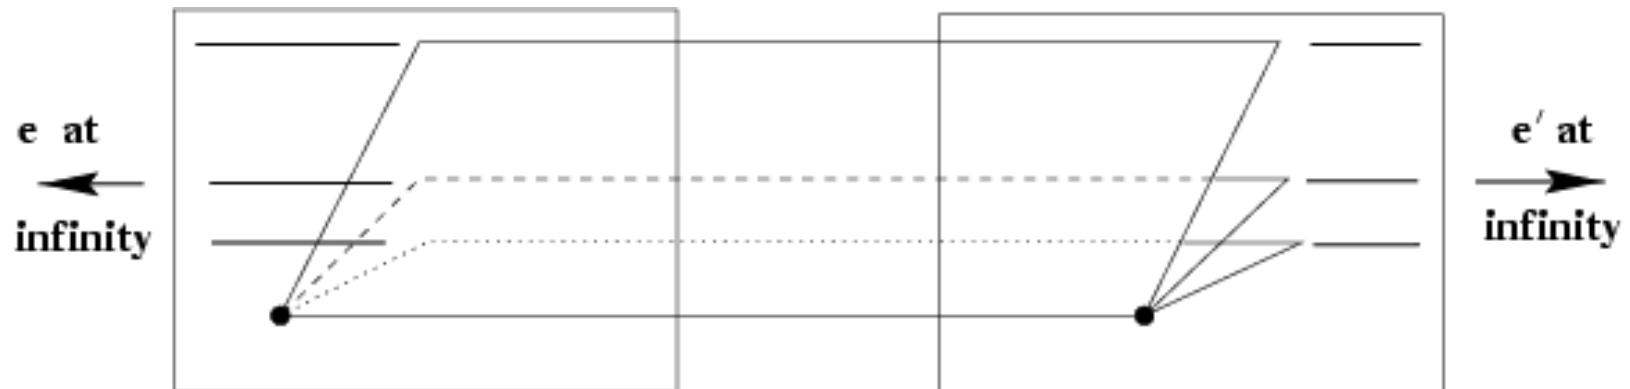
Parallel image planes



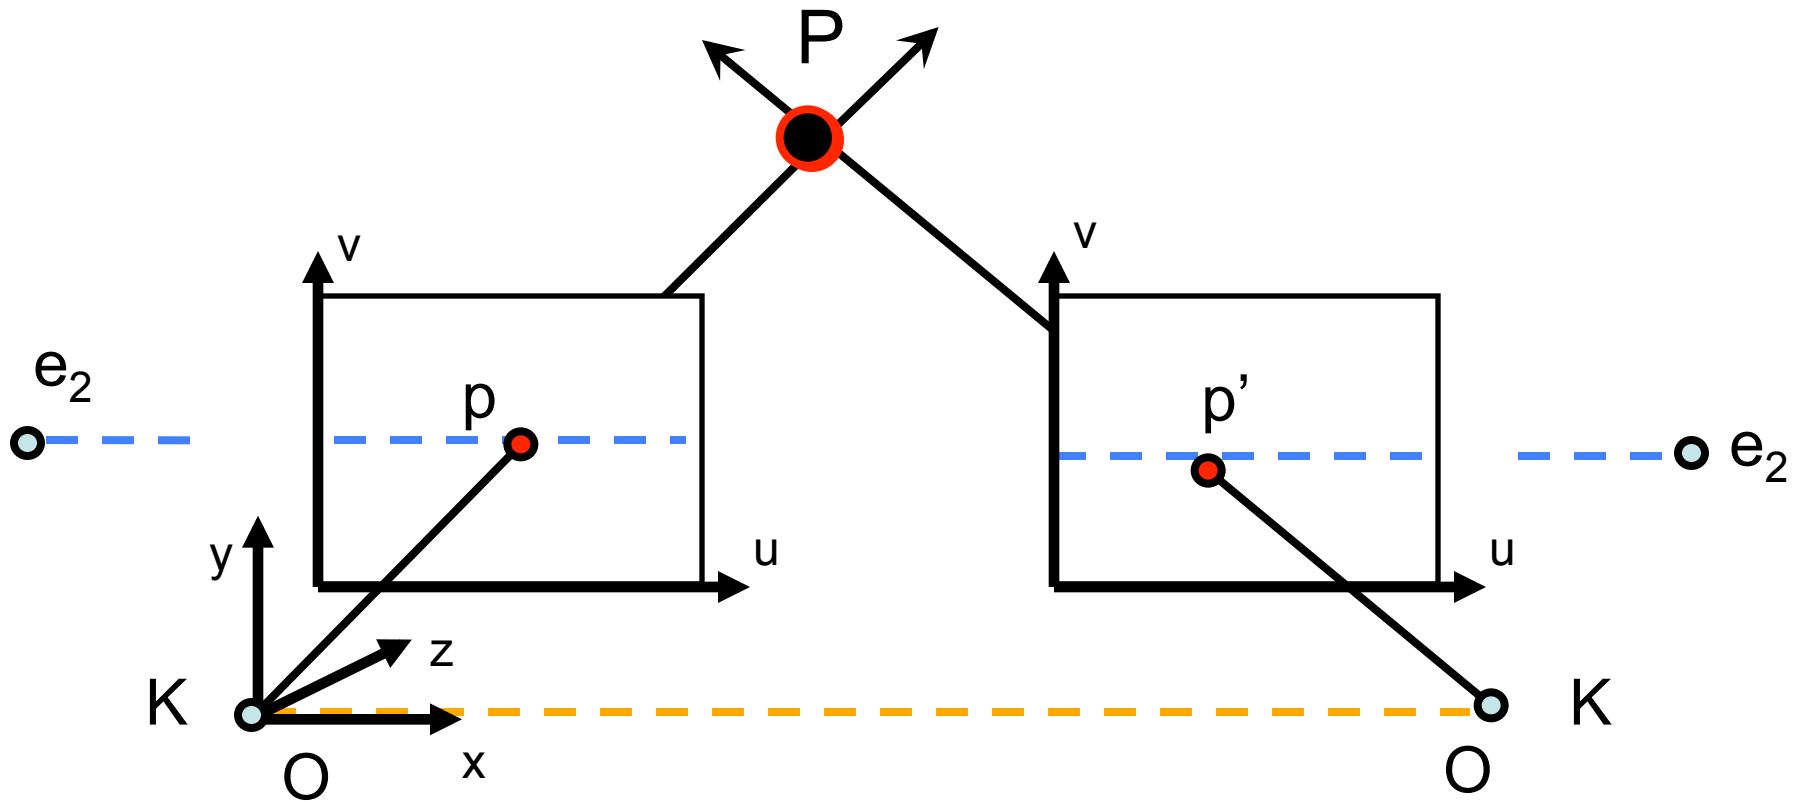
- Parallel epipolar lines
- Epipoles at infinity
- $v = v'$

Rectification: making two images “parallel”

Example: Parallel image planes



Parallel image planes



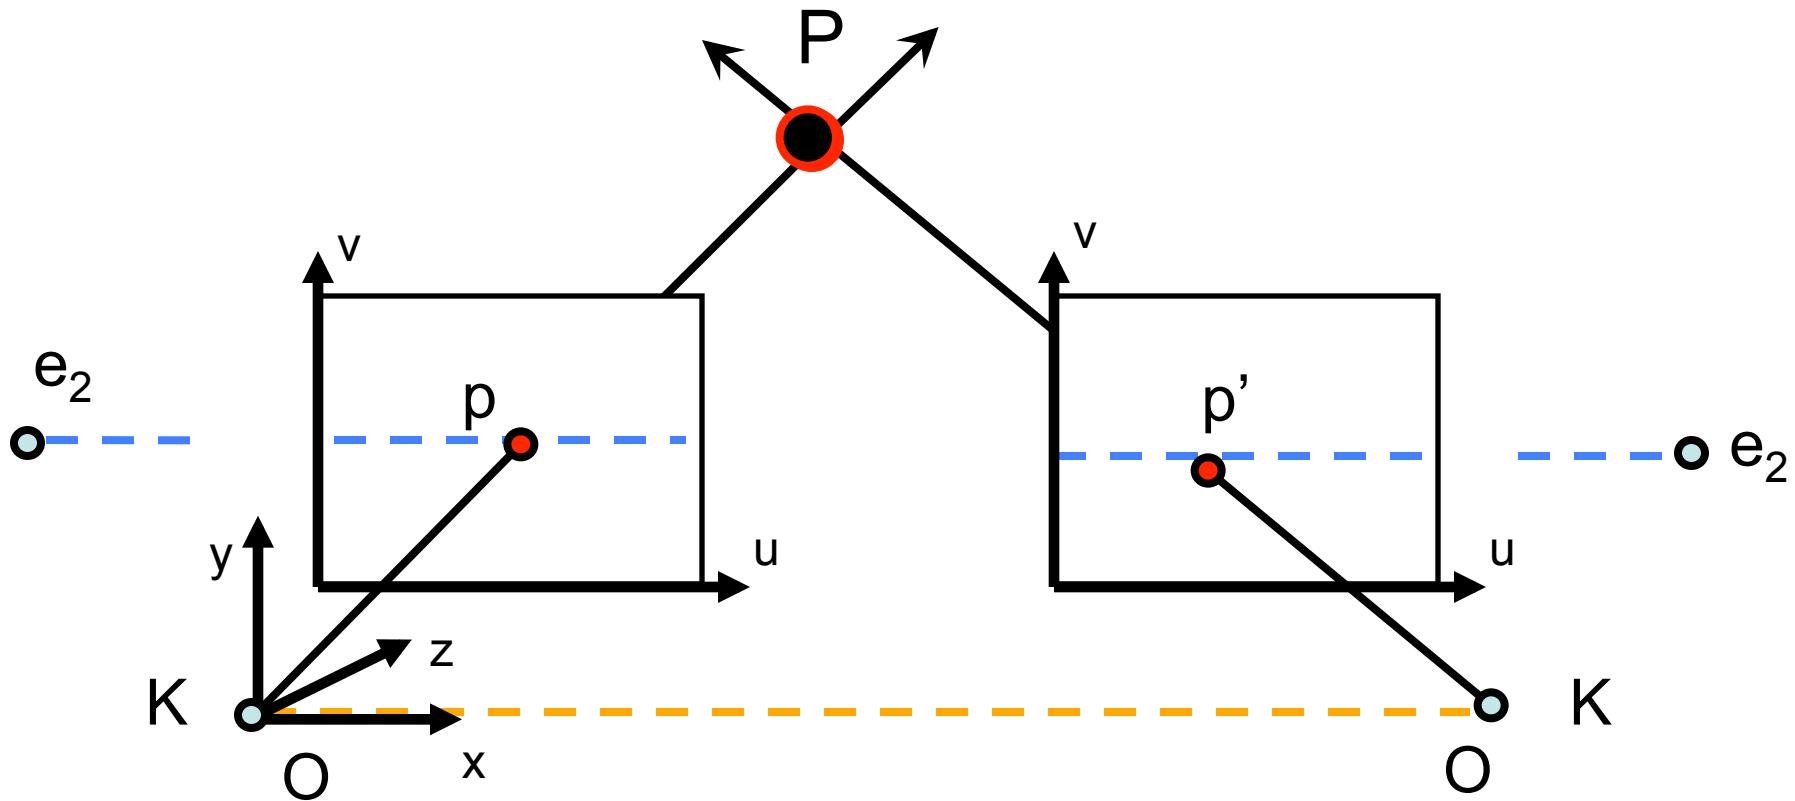
$K_1 = K_2 = \text{known}$
 $x \parallel O_1O_2$

$$E = [t_x]R$$

Cross product as matrix multiplication

$$\mathbf{a} \times \mathbf{b} = \begin{bmatrix} 0 & -a_z & a_y \\ a_z & 0 & -a_x \\ -a_y & a_x & 0 \end{bmatrix} \begin{bmatrix} b_x \\ b_y \\ b_z \end{bmatrix} = [\mathbf{a}_\times] \mathbf{b}$$

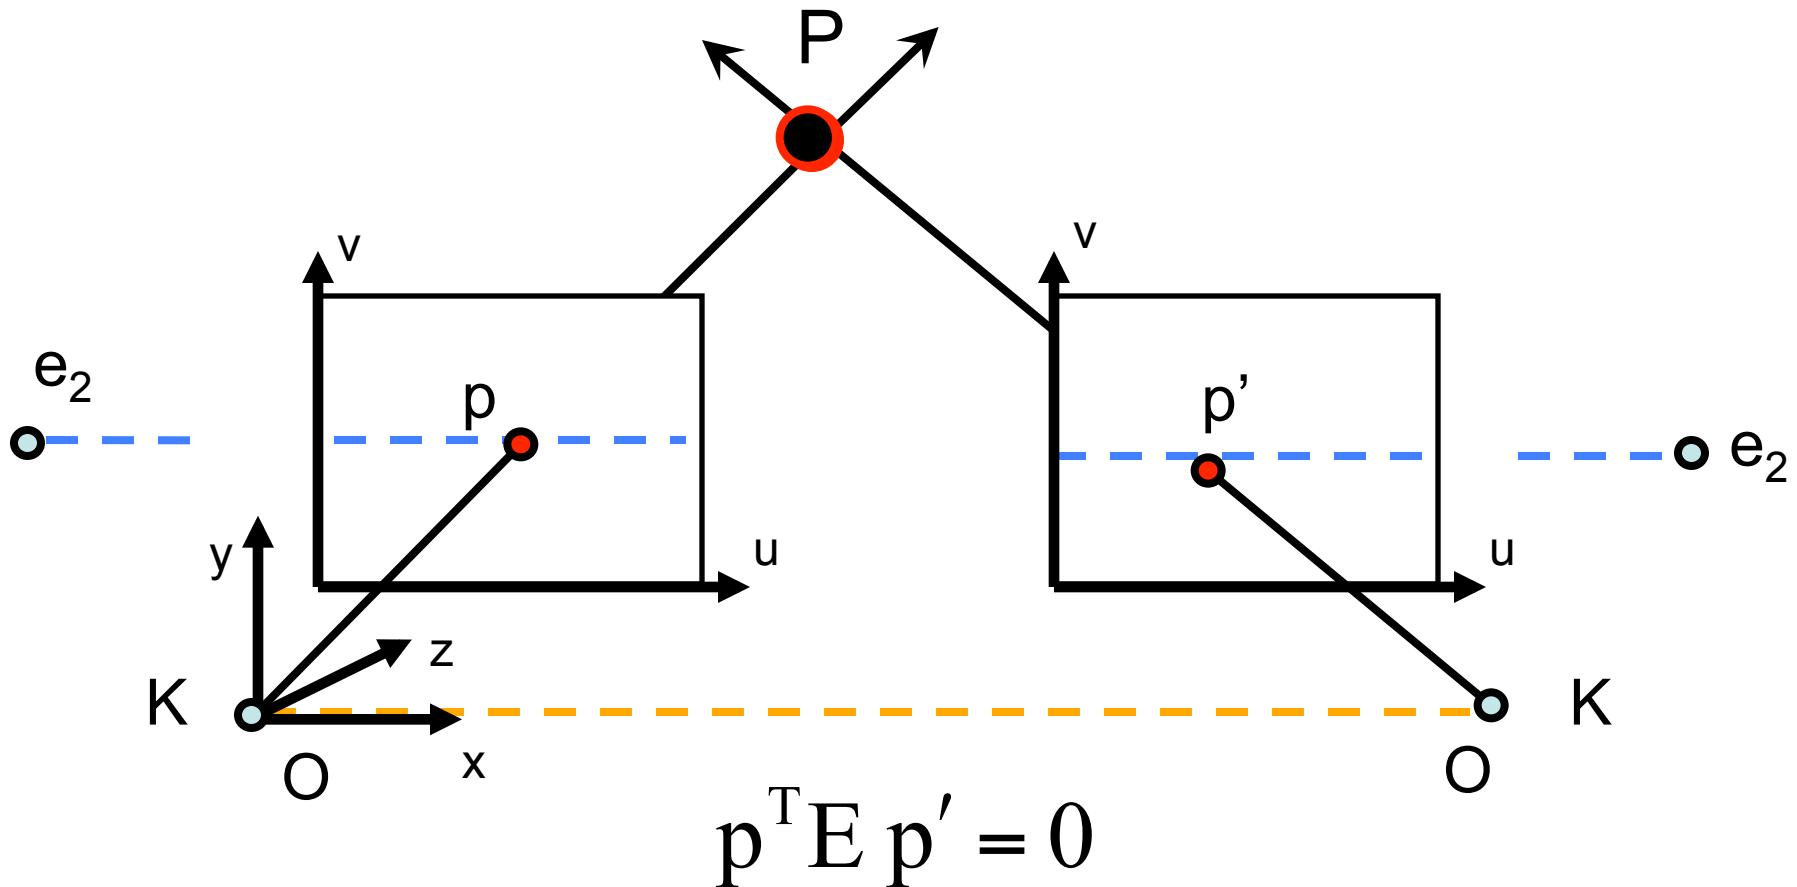
Parallel image planes



$K_1 = K_2 = \text{known}$
 $x \text{ parallel to } O_1O_2$

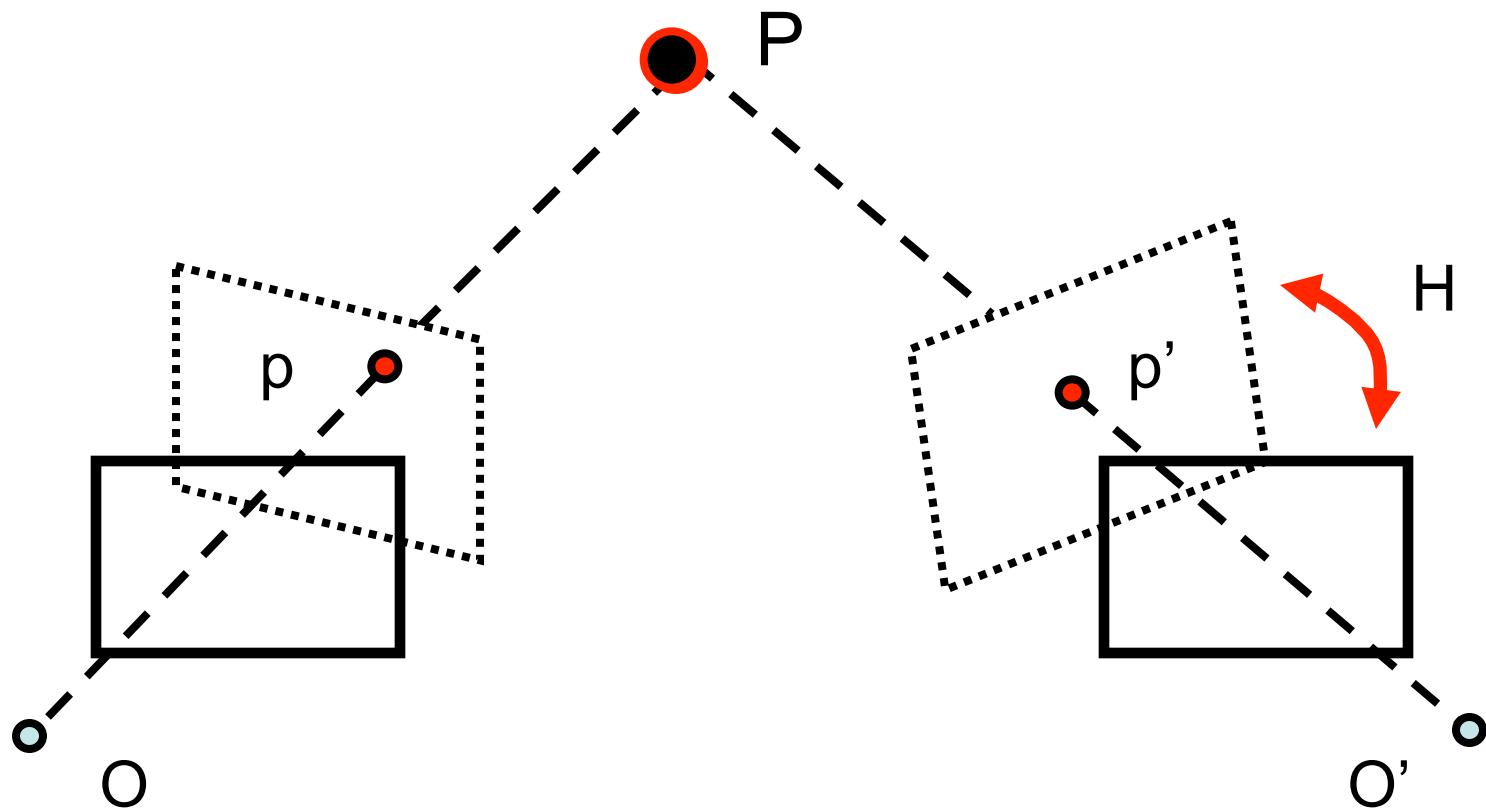
$$E = [t_x]R = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -T \\ 0 & T & 0 \end{bmatrix} \rightarrow v = v'?$$

Parallel image planes



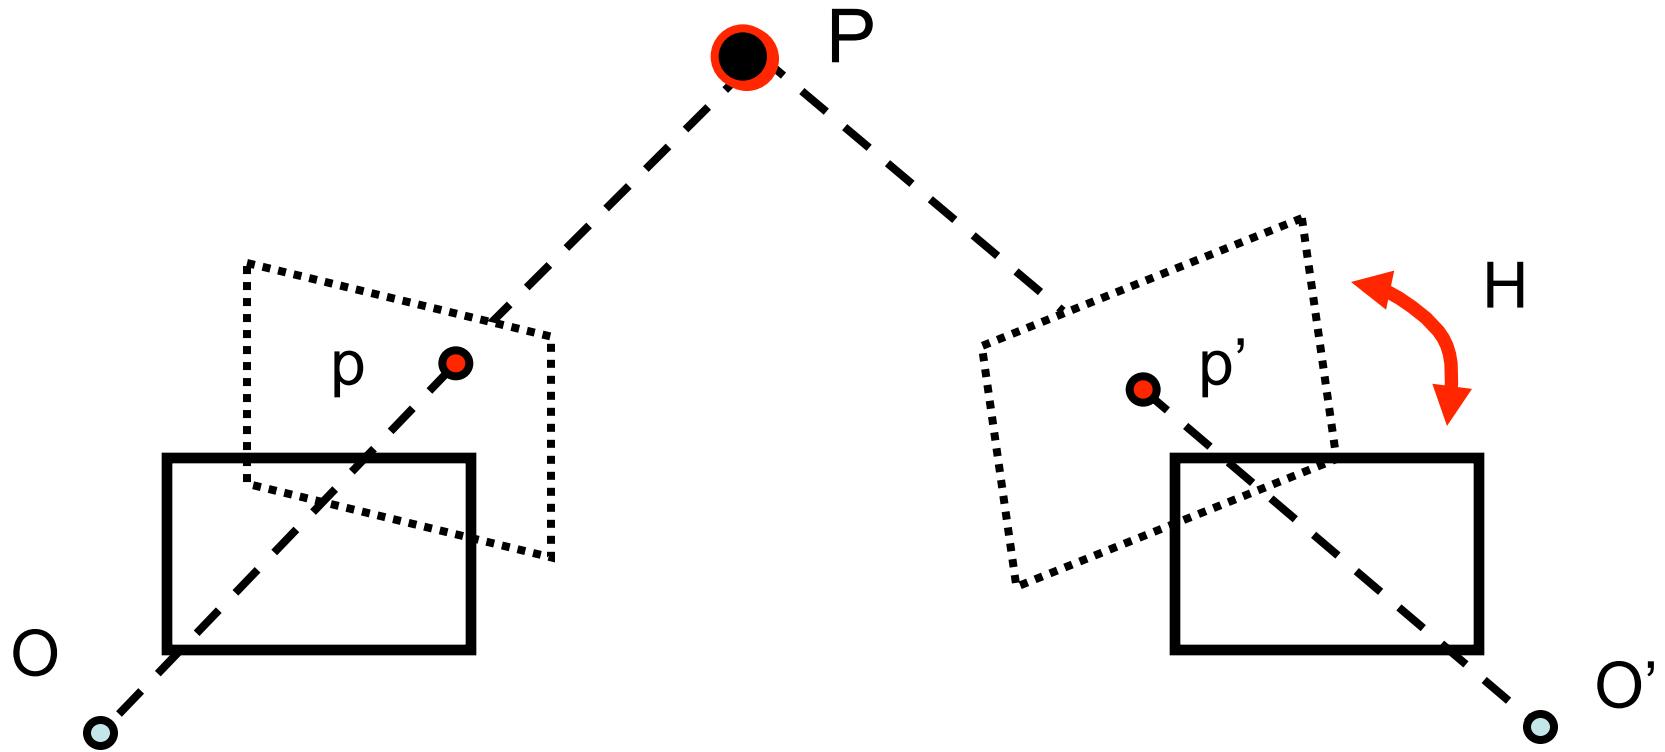
$$(u \quad v \quad 1) \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -T \\ 0 & T & 0 \end{bmatrix} \begin{pmatrix} u' \\ v' \\ 1 \end{pmatrix} = 0 \quad (u \quad v \quad 1) \begin{pmatrix} 0 \\ -T \\ Tv' \end{pmatrix} = 0 \quad Tv = Tv'$$

Making image planes parallel



GOAL: Estimate the perspective transformation H
that makes the images parallel

Making image planes parallel

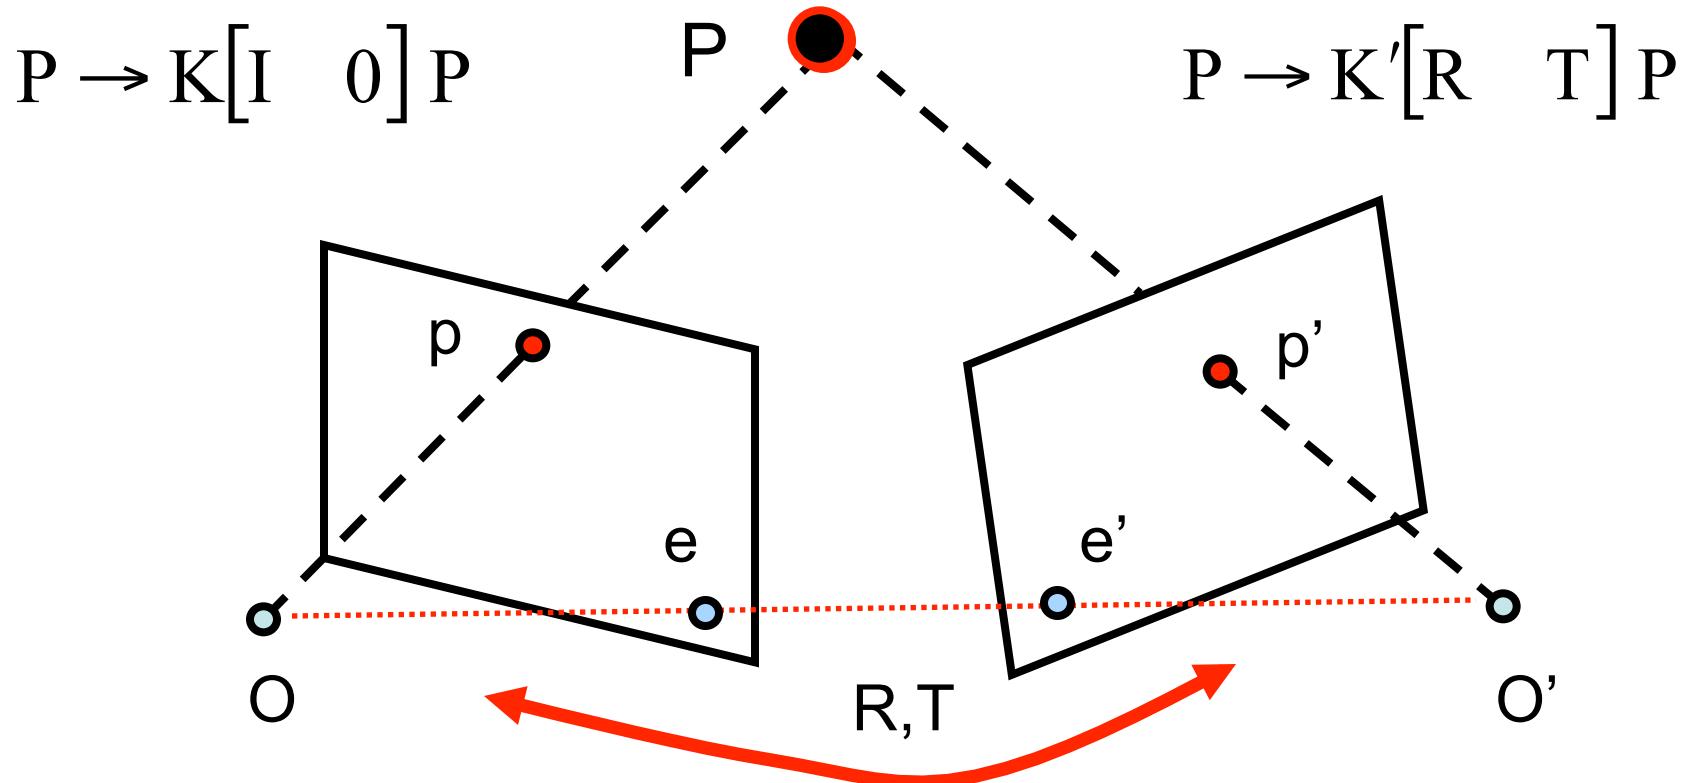


GOAL: Estimate the perspective transformation H that makes images parallel

Impose $v' = v$

- This leaves degrees of freedom for determining H
- If not appropriate H is chosen, severe projective distortions on image take place
- We impose a number of restriction while computing H

Making image planes parallel

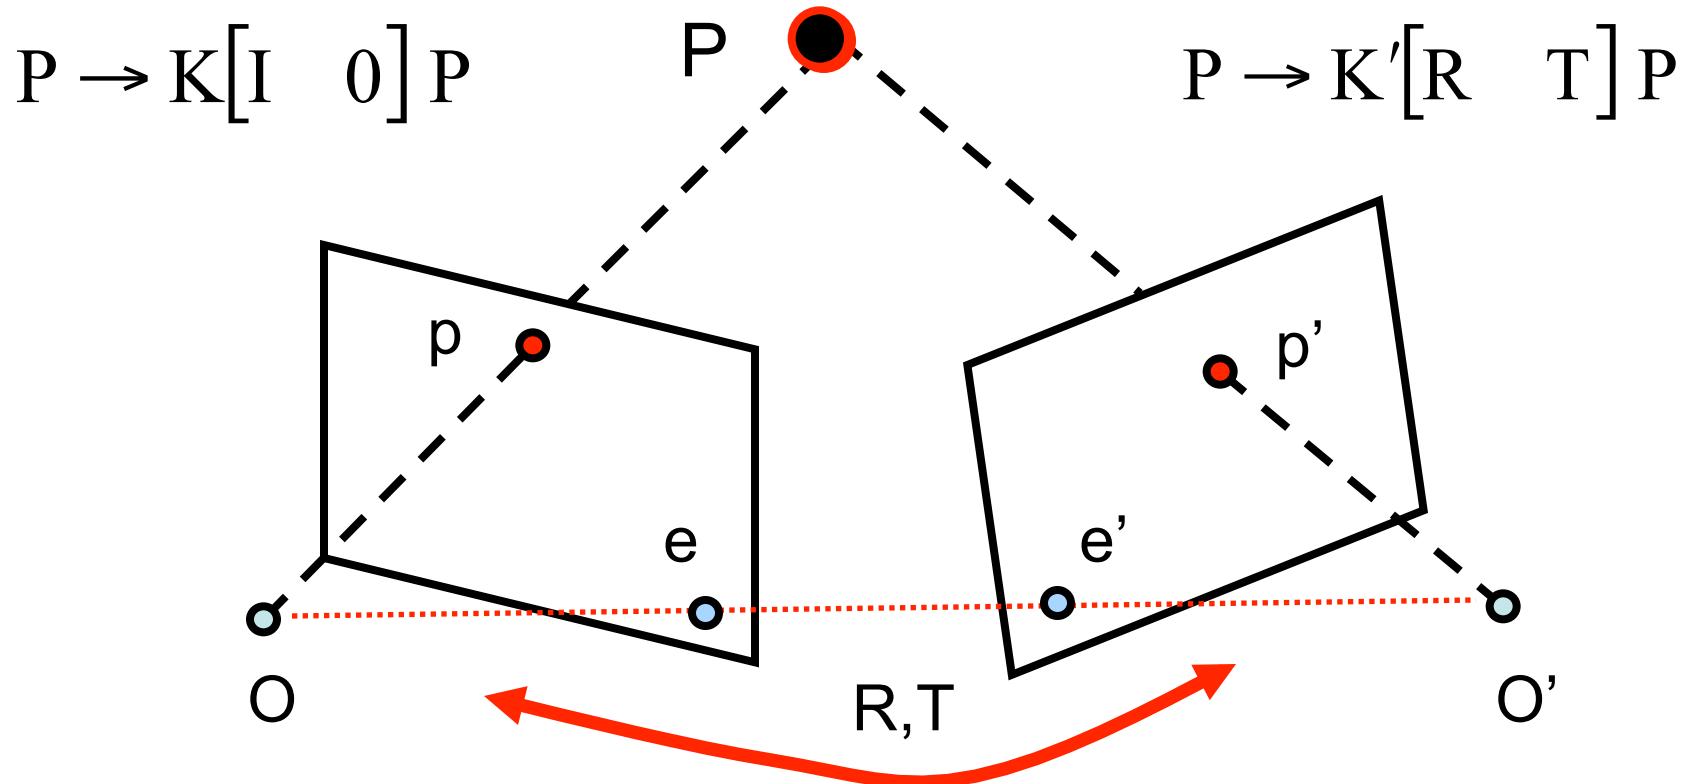


0. Compute epipoles

$$e = [\ e_1 \quad e_2 \quad 1 \]^T = -K R^T T$$

$$e' = K' T$$

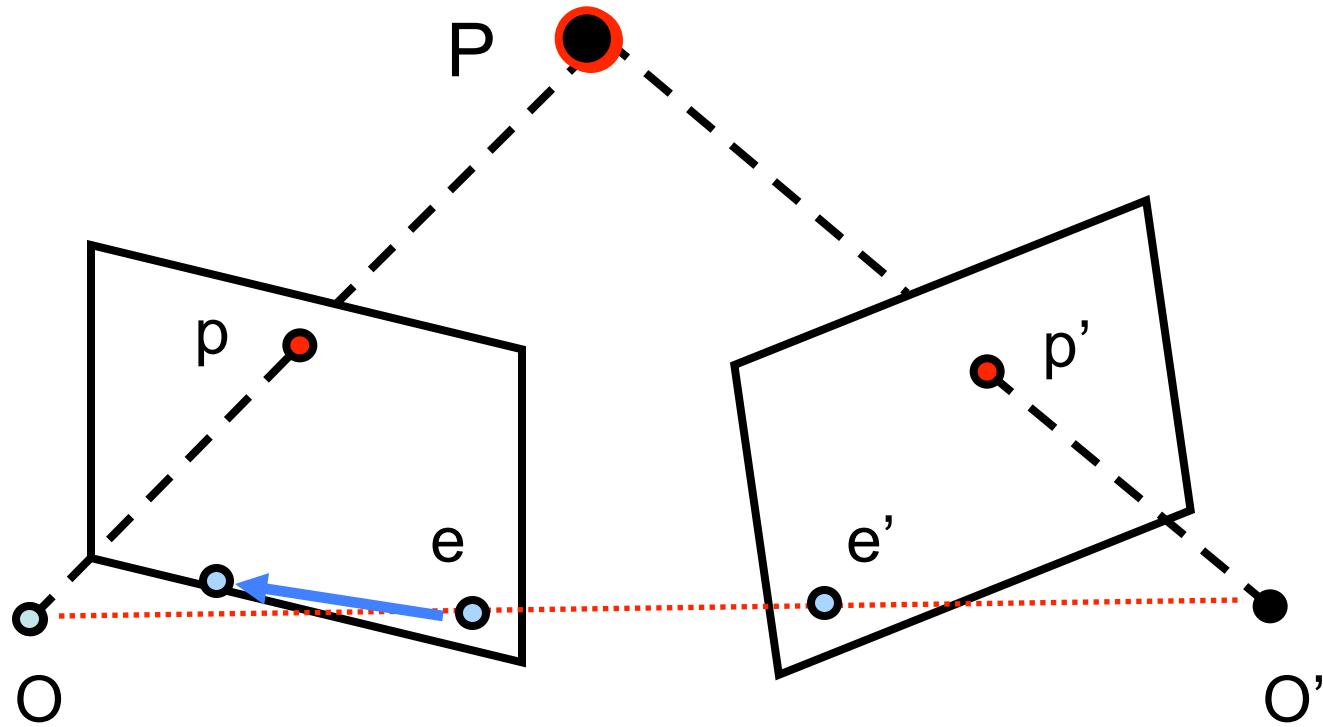
Making image planes parallel



0. Compute epipoles

$$K \begin{bmatrix} I & 0 \end{bmatrix} \begin{bmatrix} -R^T T \\ 1 \end{bmatrix} = -KR^T T \quad e' = K' \begin{bmatrix} R & T \end{bmatrix} \begin{bmatrix} 0 \\ 1 \end{bmatrix} = K'T$$

Making image planes parallel

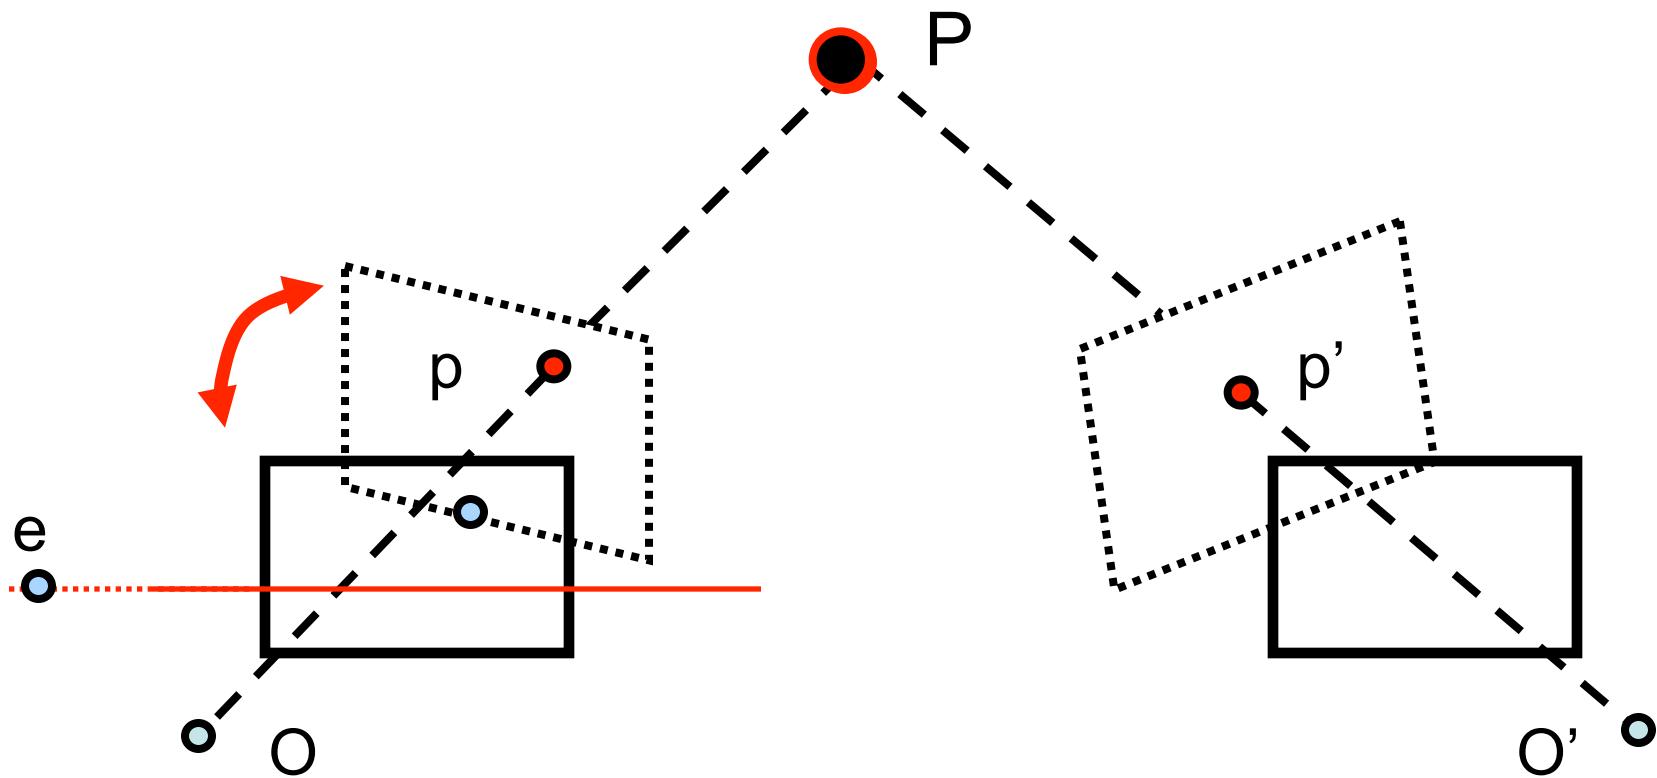


1. Map e to the x -axis at location $[1, 0, 1]^T$ (normalization)

$$e = [e_1 \ e_2 \ 1]^T \rightarrow [1 \ 0 \ 1]^T$$

$$H_1 = R_H T_H$$

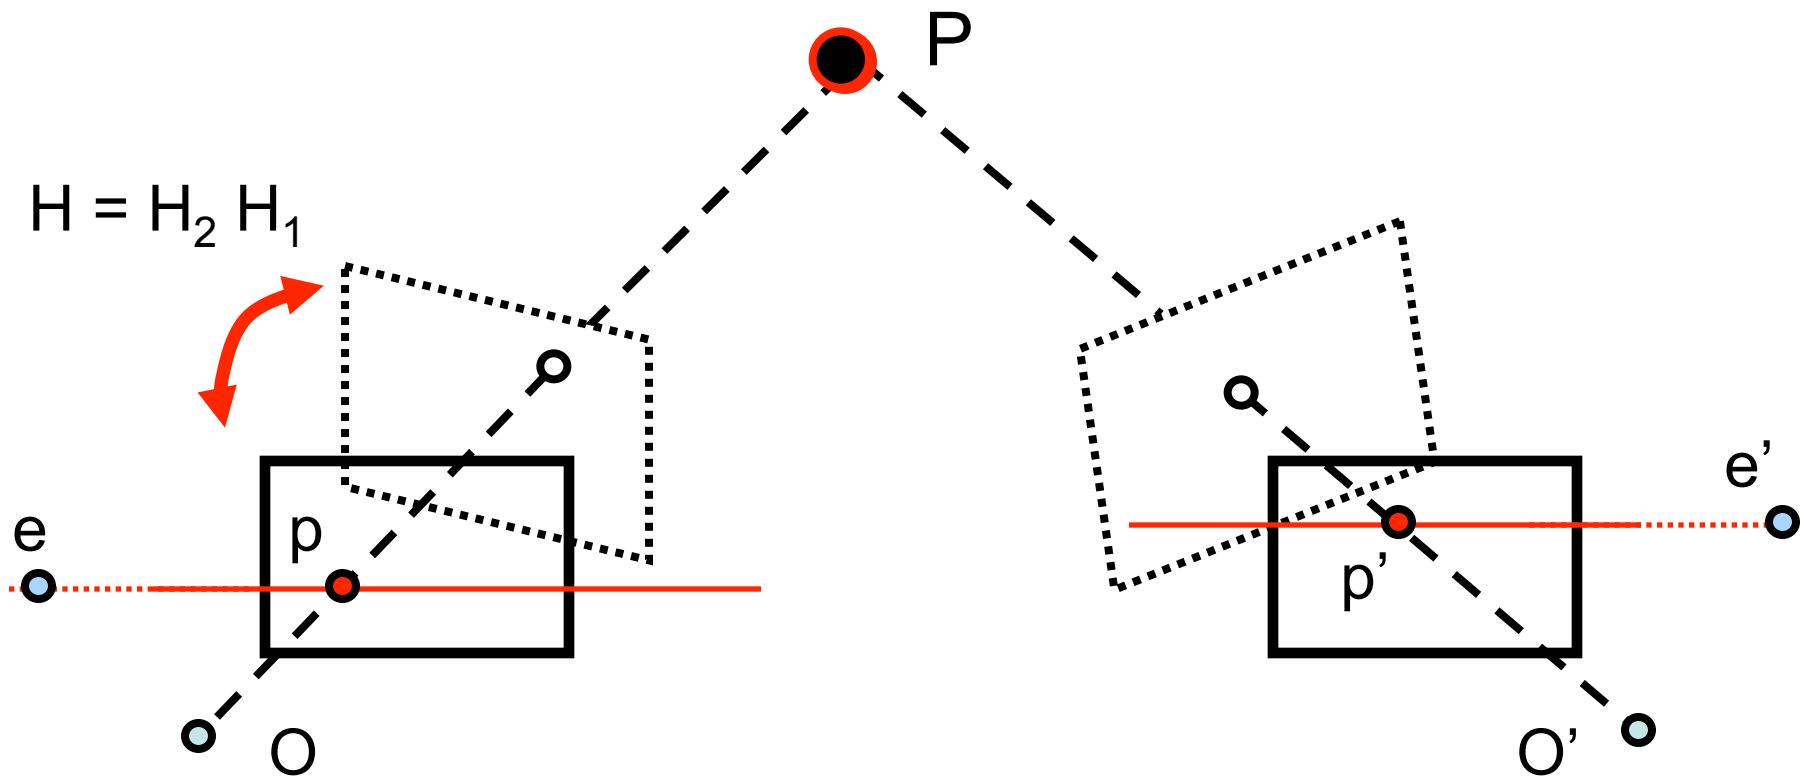
Making image planes parallel



2. Send epipole to infinity: $e = [1 \ 0 \ 1]^T \rightarrow [1 \ 0 \ 0]^T$

$$H_2 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$

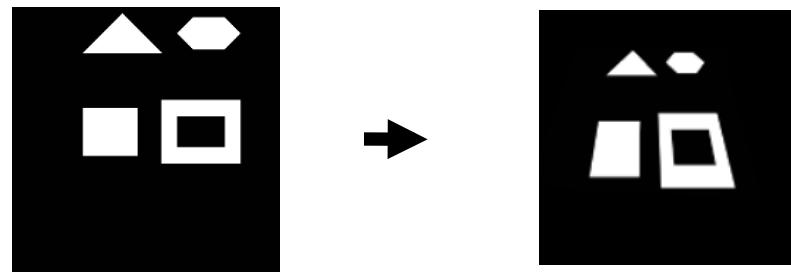
Making image planes parallel



3. Define: $H = H_2 H_1$
4. Align epipolar lines

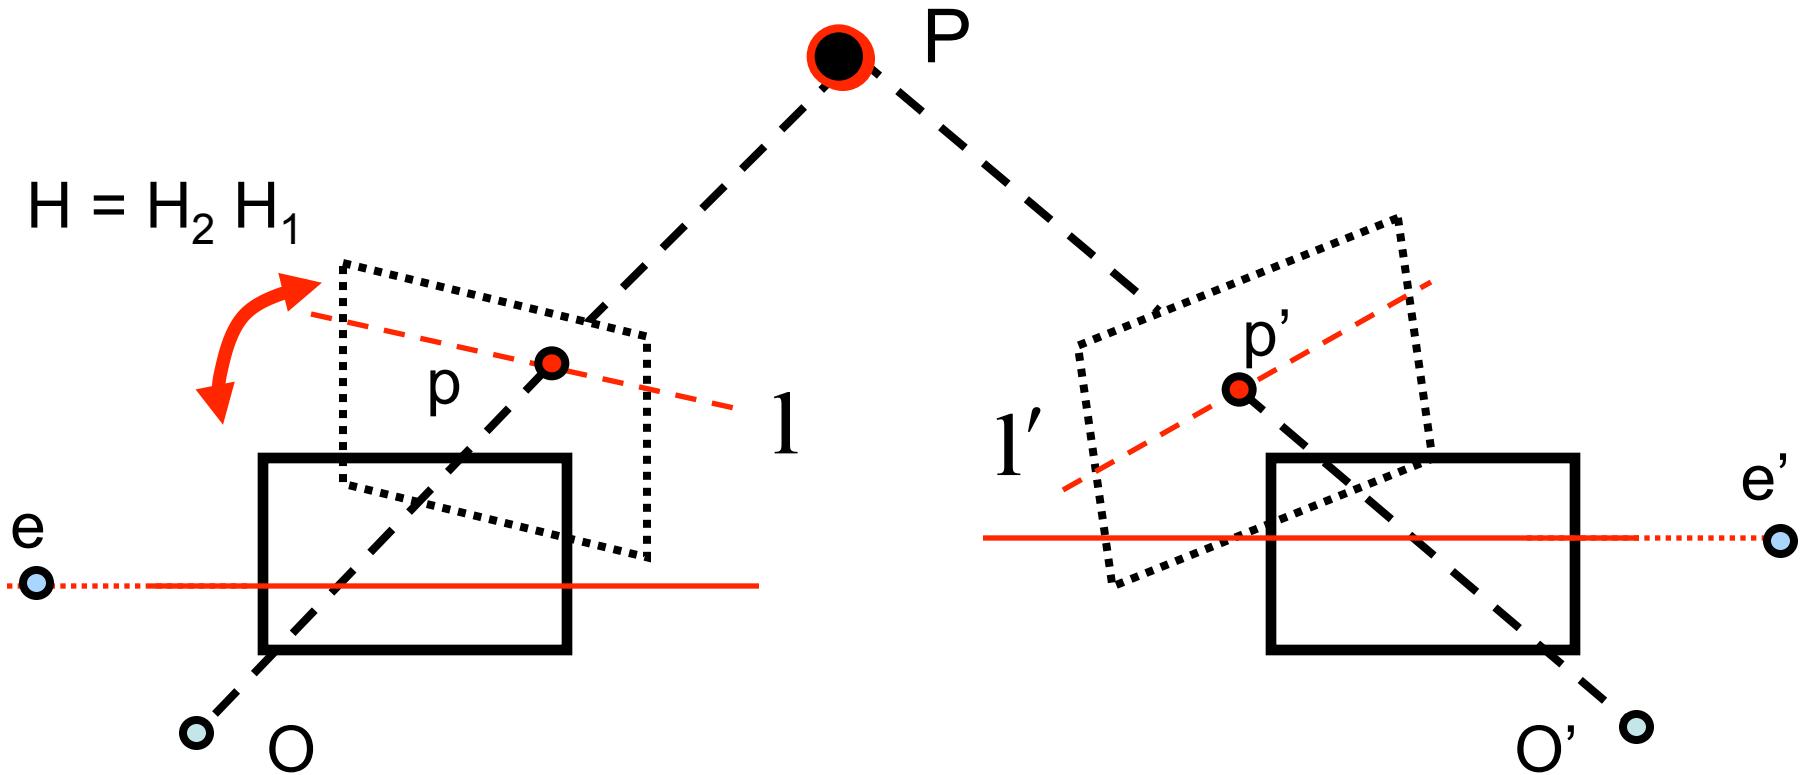
Projective transformation of a line (in 2D)

$$H = \begin{bmatrix} A & t \\ v & b \end{bmatrix}$$



$$l' \rightarrow H^{-T} l$$

Making image planes parallel

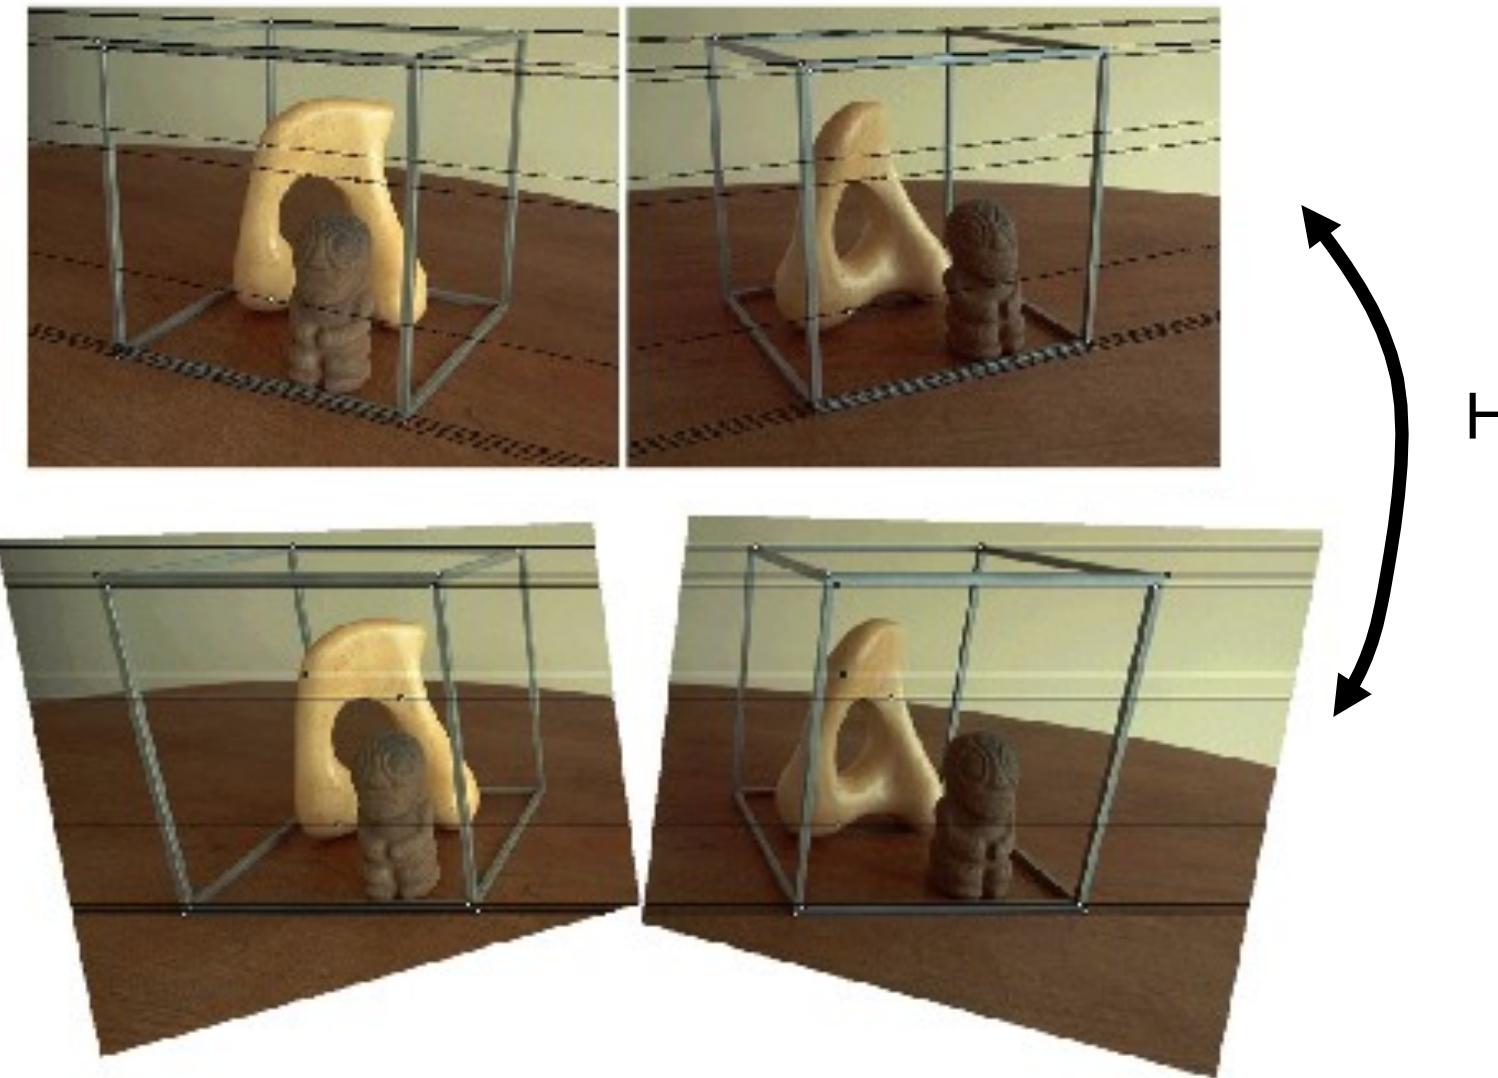


3. Define: $H = H_2 H_1$
4. Align epipolar lines

$$\overline{H}'^{-T} l' = \overline{H}^{-T} l$$

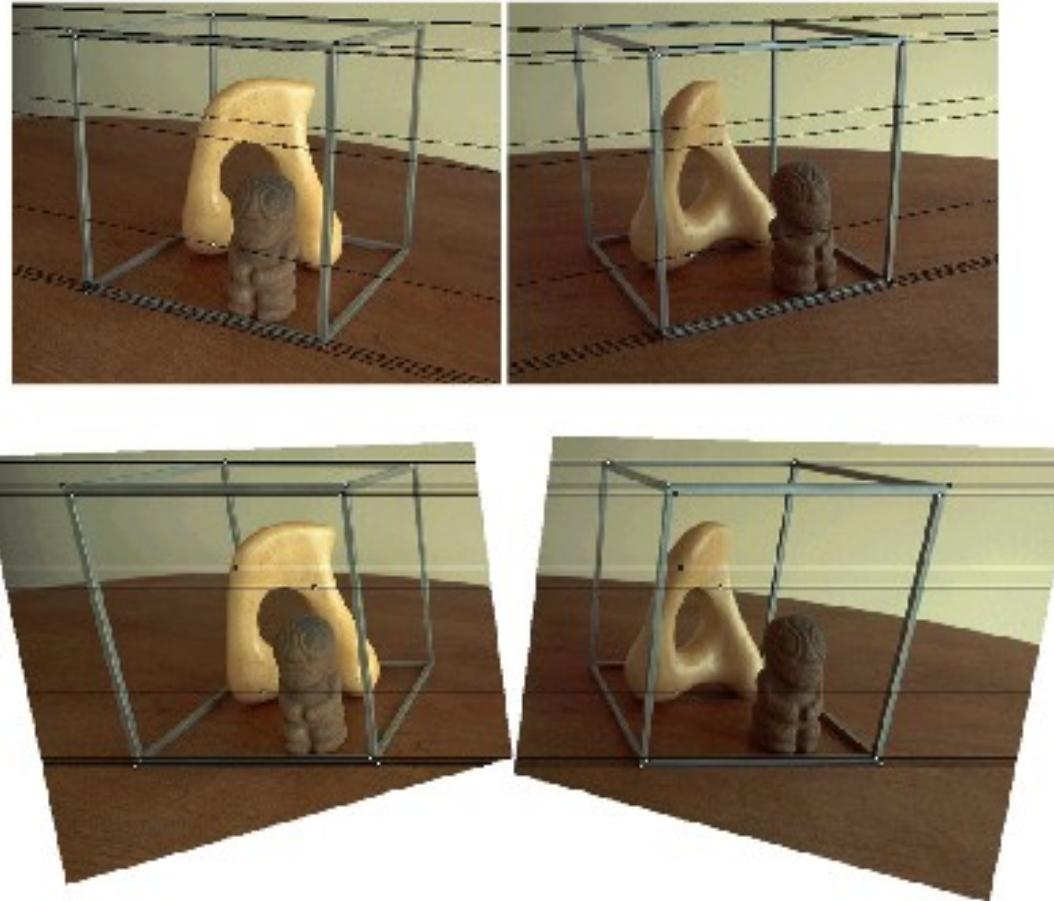
These are called **matched pair** of transformation

Making image planes parallel



Courtesy figure S. Lazebnik

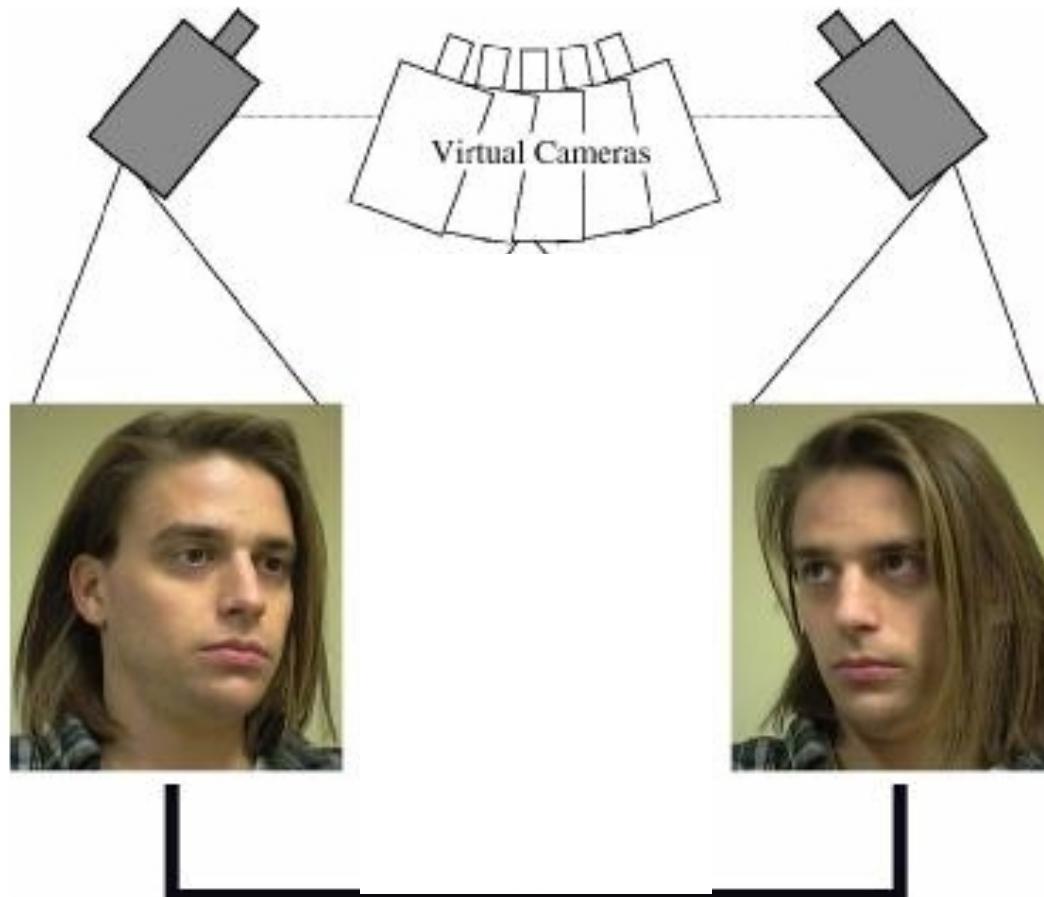
Why rectification is useful?



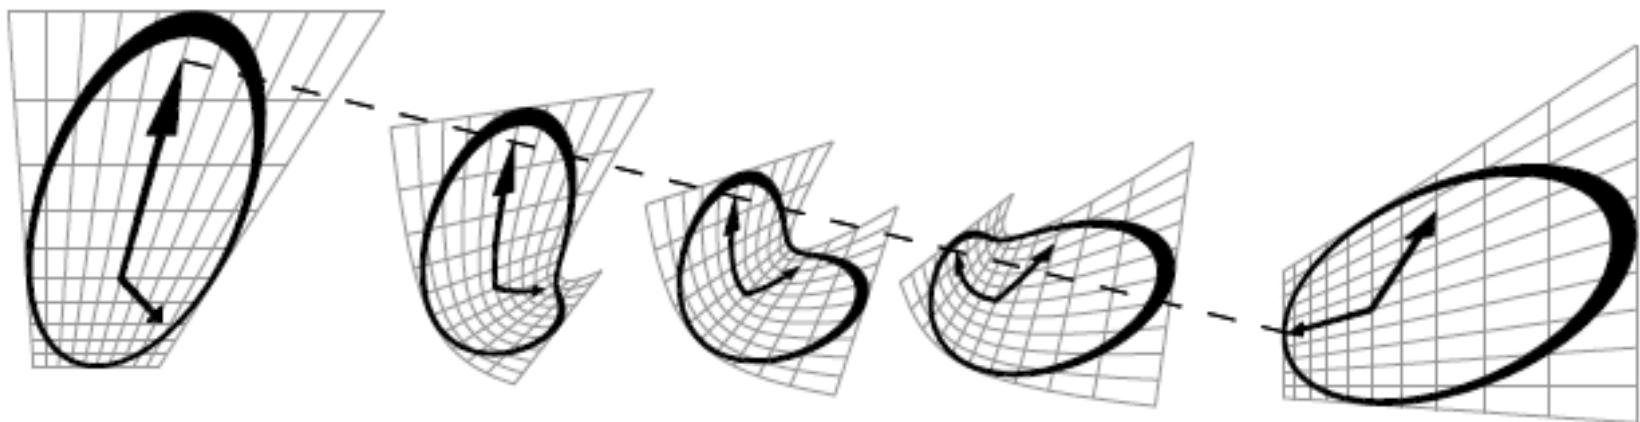
- Makes the correspondence problem easier
- Makes triangulation easy

Application: view morphing

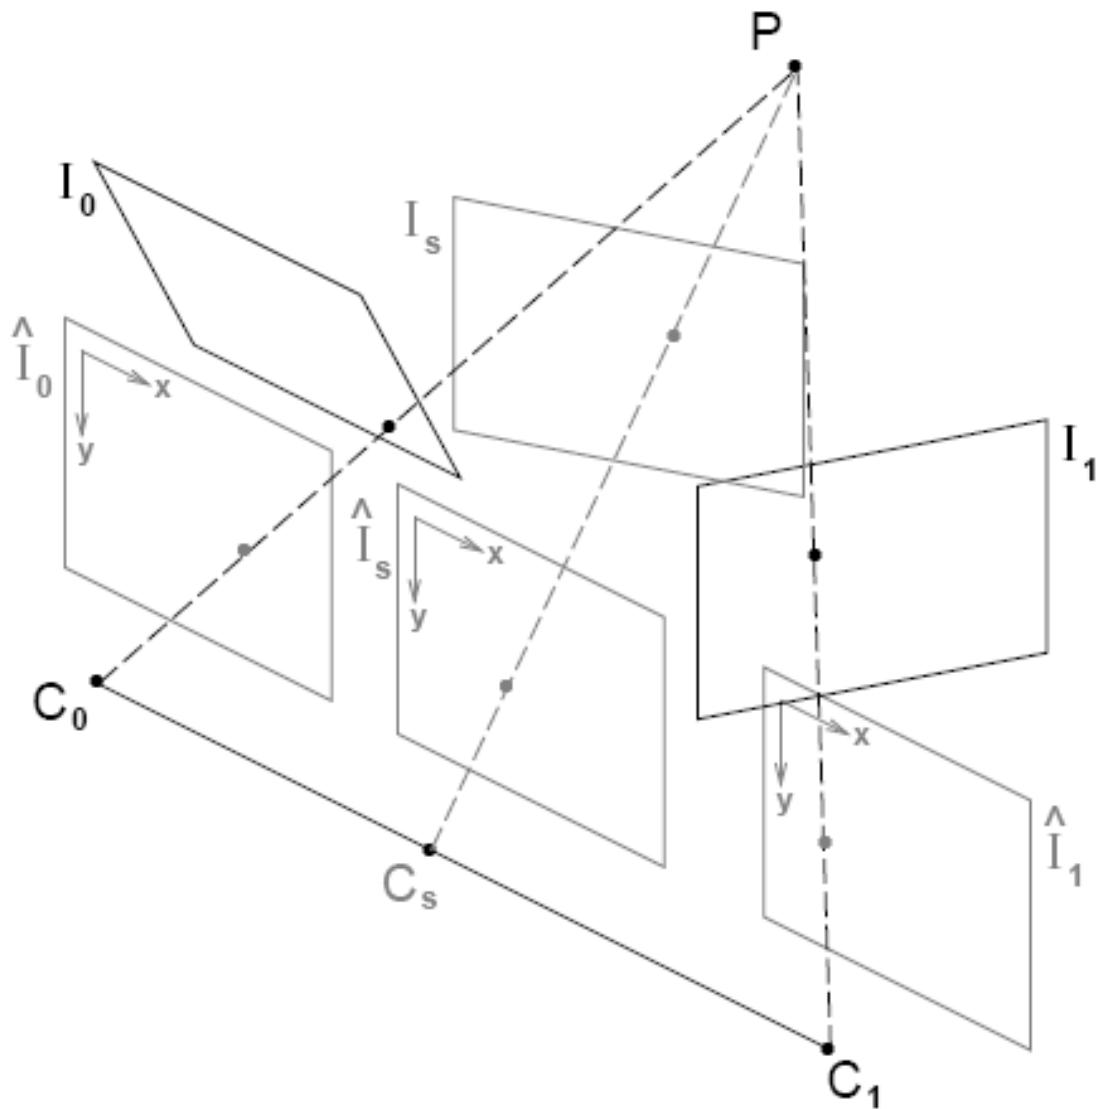
S. M. Seitz and C. R. Dyer, *Proc. SIGGRAPH 96*, 1996, 21-30



Morphing without using geometry

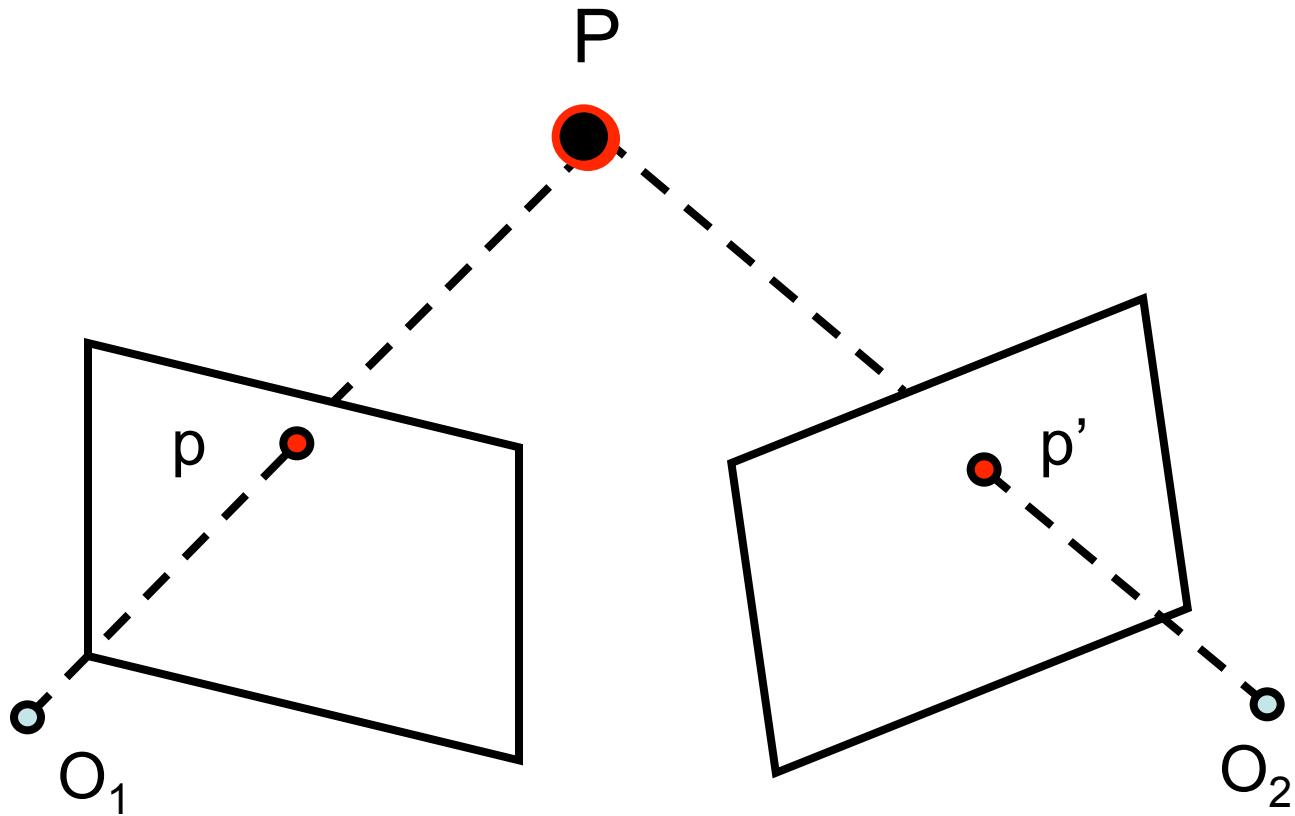


Rectification





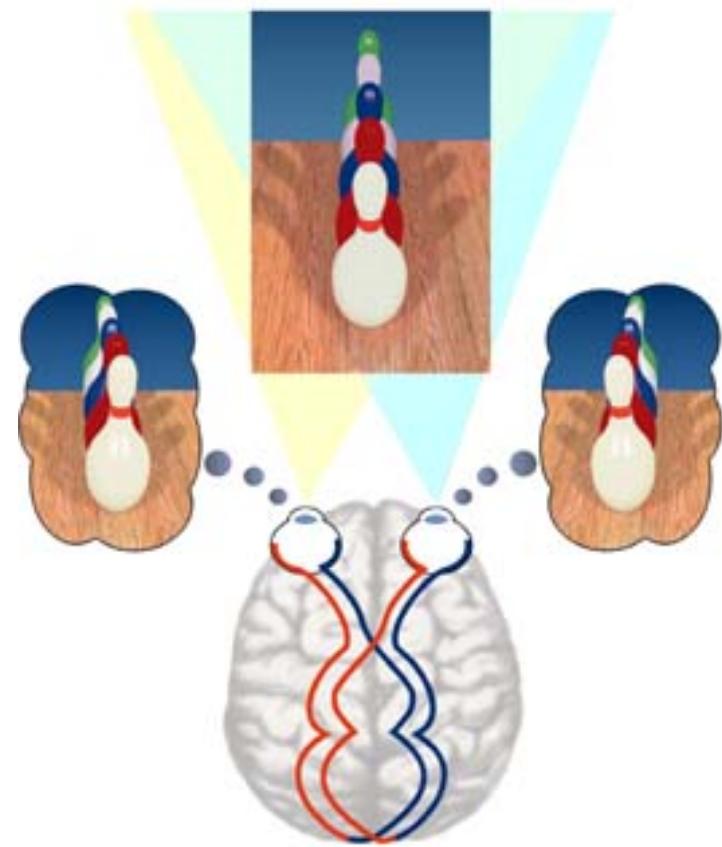
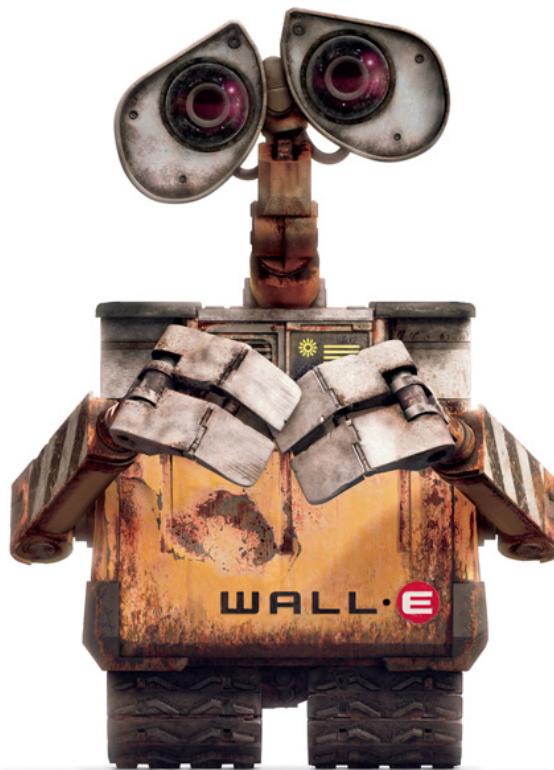
Stereo vision



Subgoals:

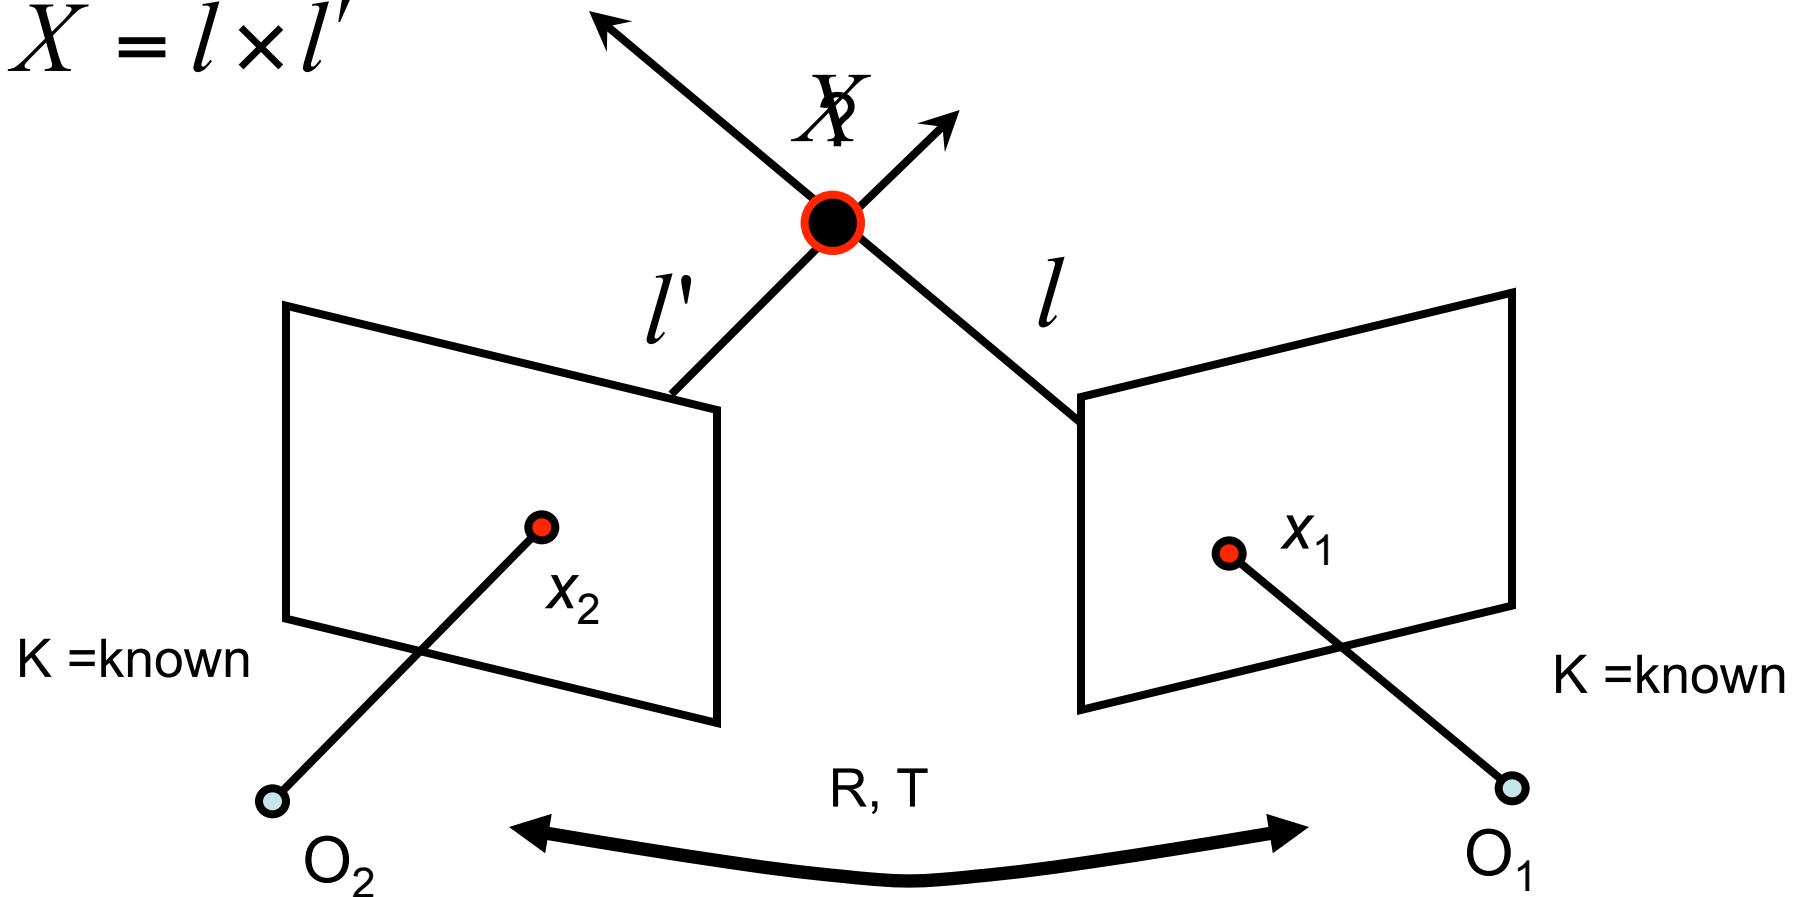
- Solve the correspondence problem
- Use corresponding observations to triangulate

Two eyes help!



Two eyes help!

$$X = l \times l'$$

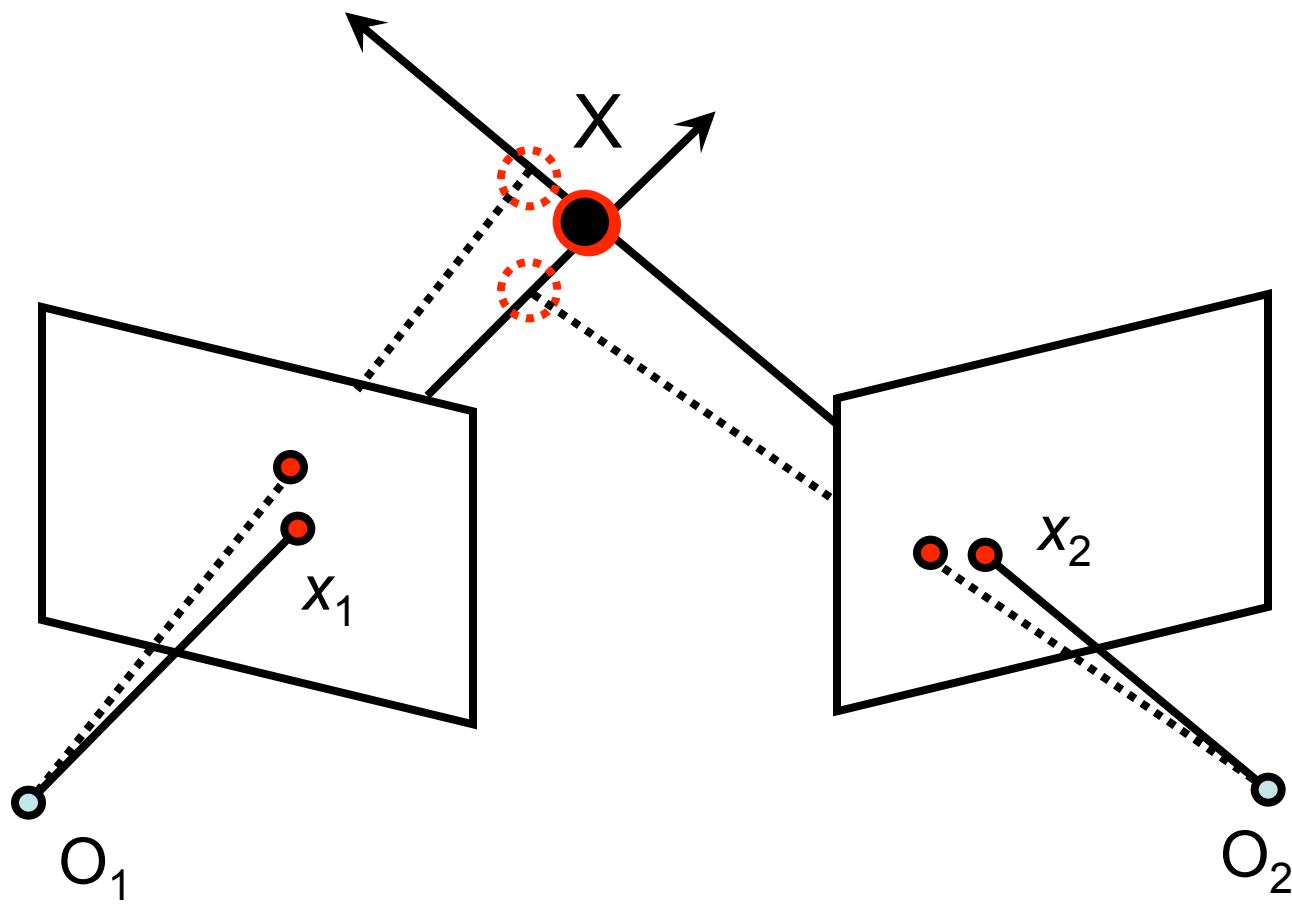


This is called **triangulation**

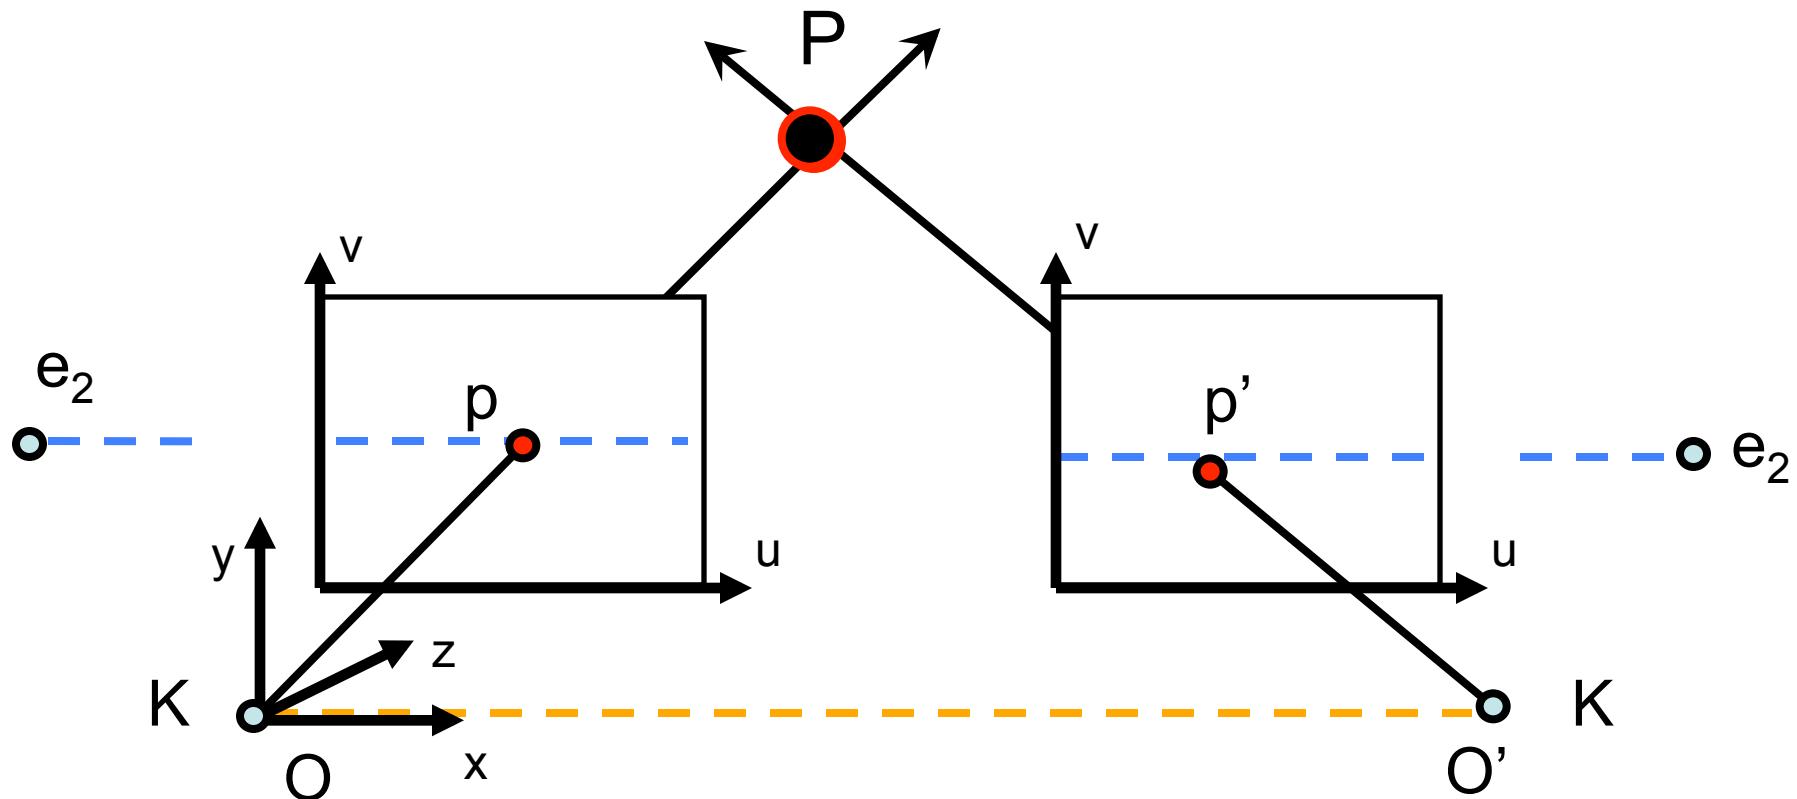
Triangulation

- Find X that minimizes

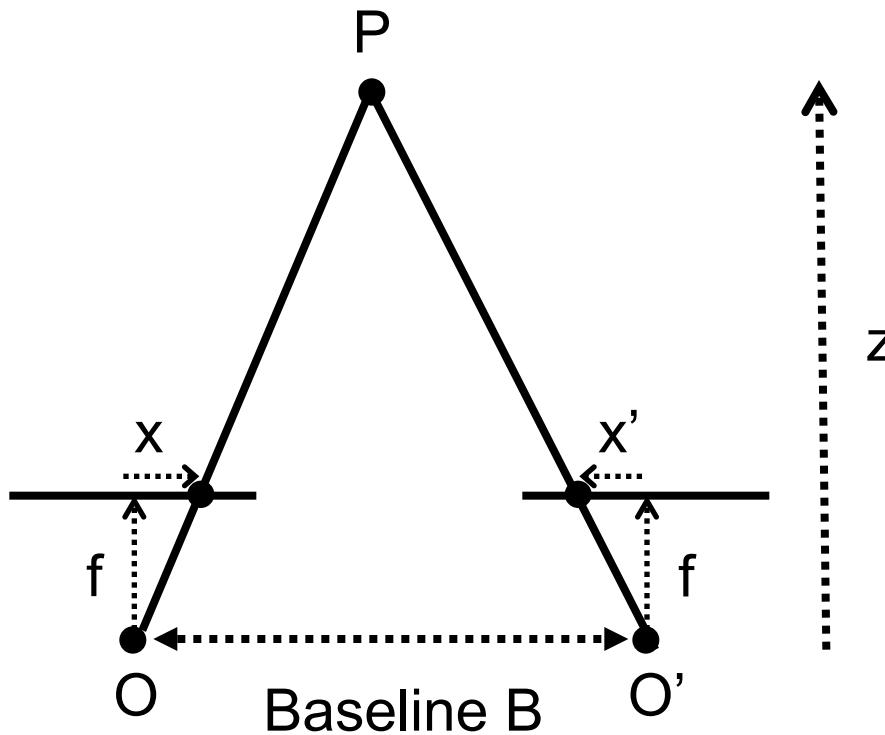
$$d^2(x_1, M_1 X) + d^2(x_2, M_2 X)$$



Computing depth



Computing depth



$$x - x' = \frac{B \cdot f}{z} = \text{disparity}$$

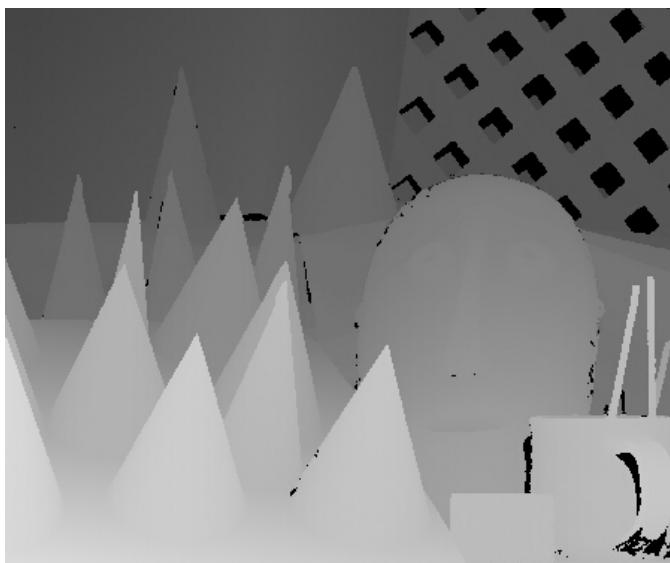
Note: Disparity is inversely proportional to depth

Disparity maps

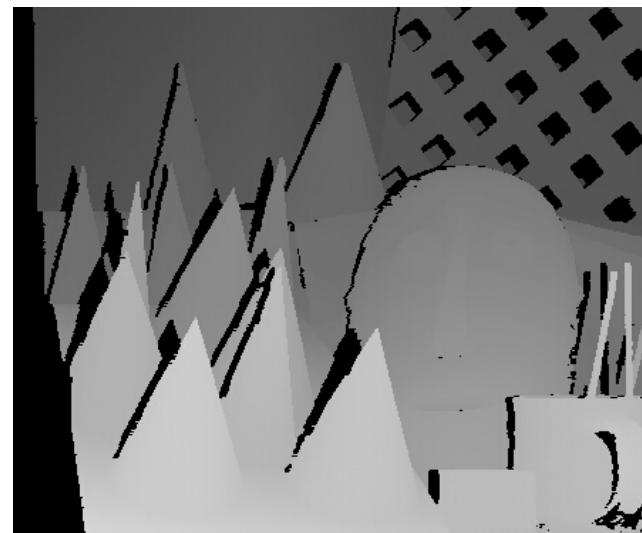


$$x - x' = \frac{B \cdot f}{z}$$

Stereo pair

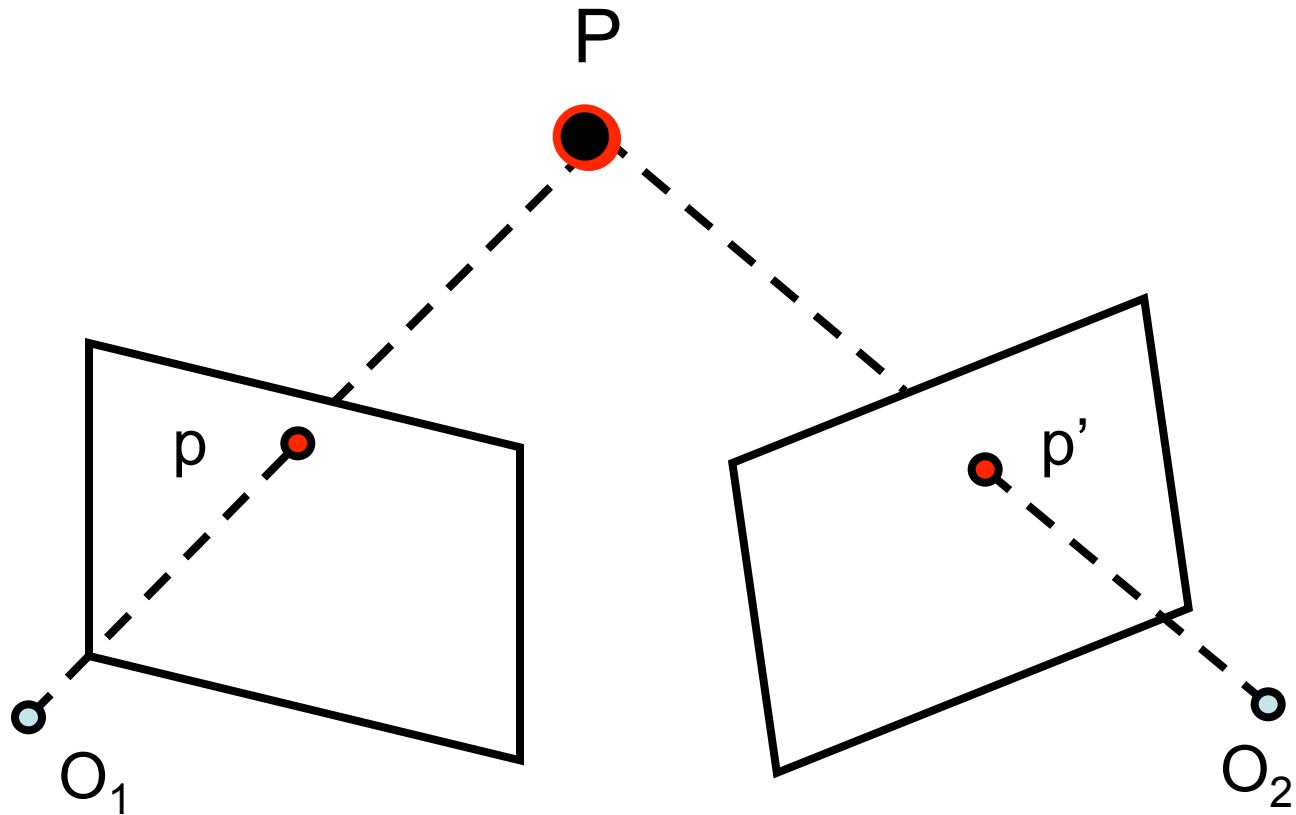


Disparity map / depth map



Disparity map with occlusions

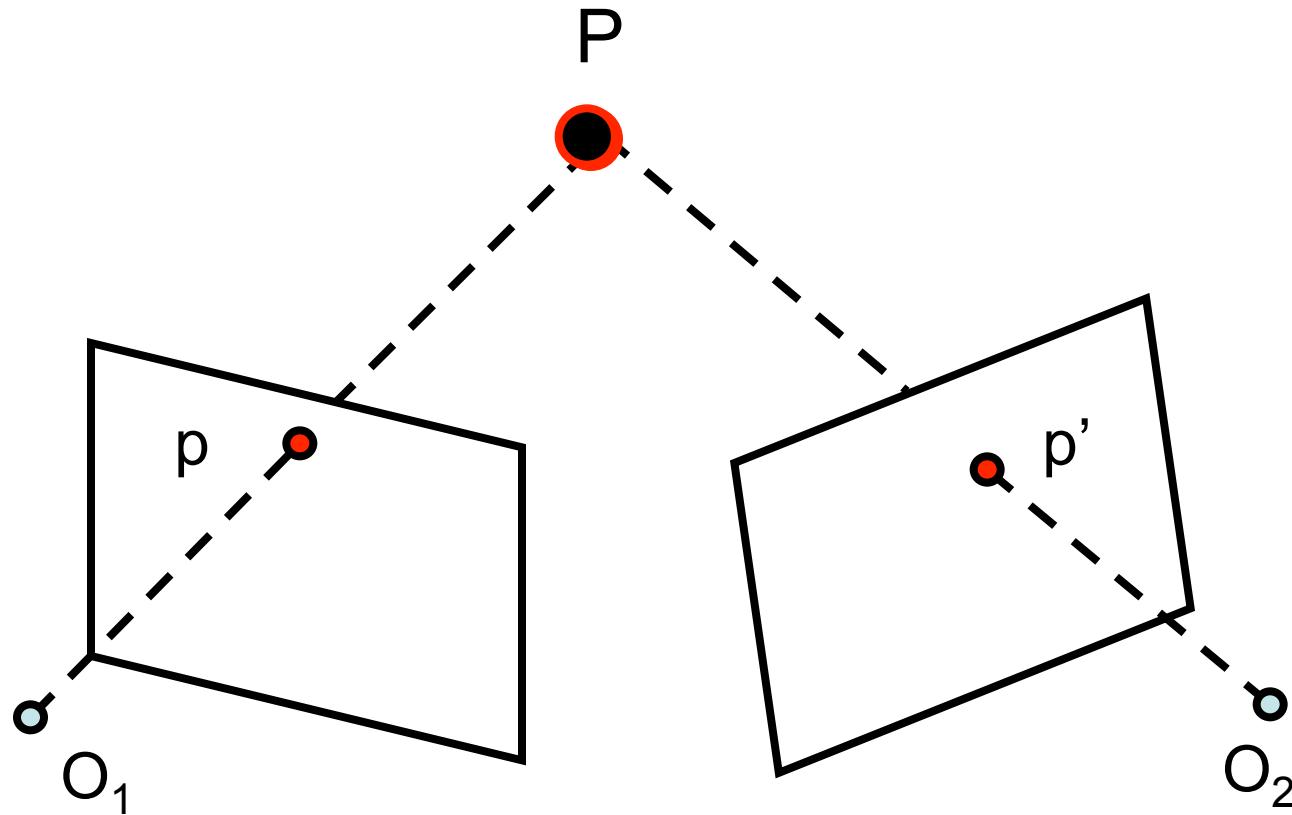
Stereo vision



Subgoals:

- Solve the correspondence problem
- Use corresponding observations to triangulate

Correspondence problem



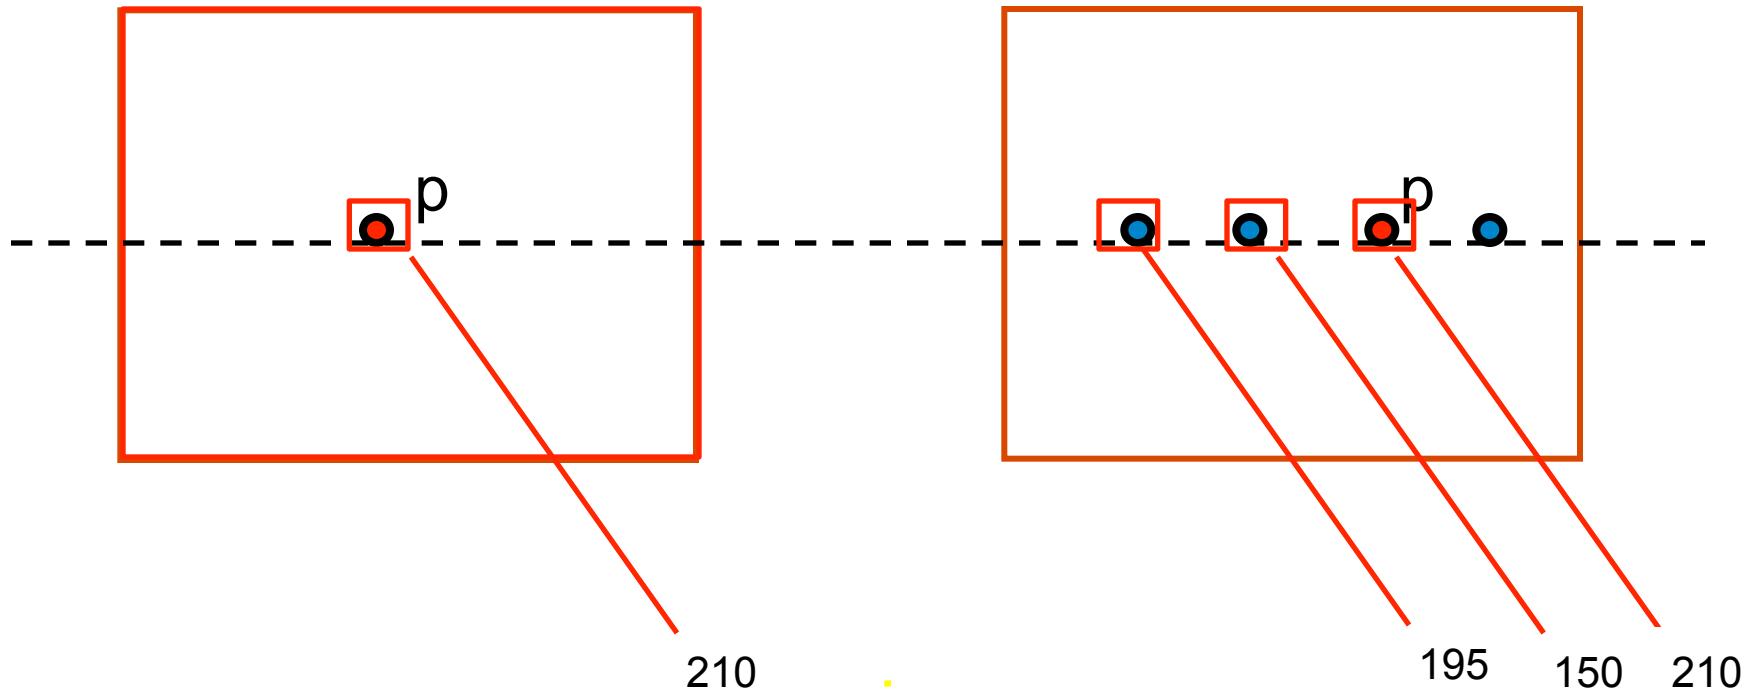
Given a point in 3d, discover corresponding observations
in left and right images [also called binocular fusion problem]

Correspondence problem

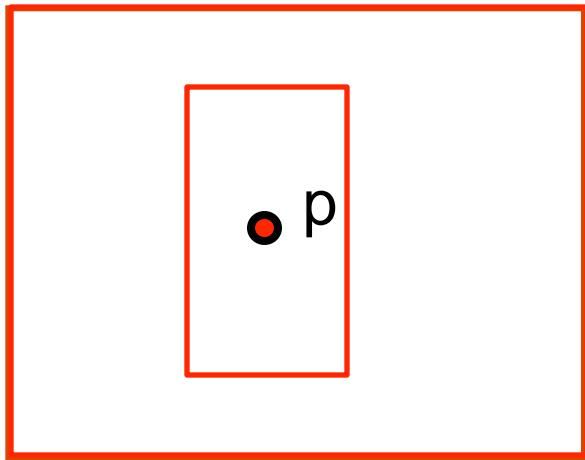
- A Cooperative Model (Marr and Poggio, 1976)
- Correlation Methods (1970--)
- Multi-Scale Edge Matching (Marr, Poggio and Grimson, 1979-81)

[FP] Chapters: 11

Correlation Methods (1970--)

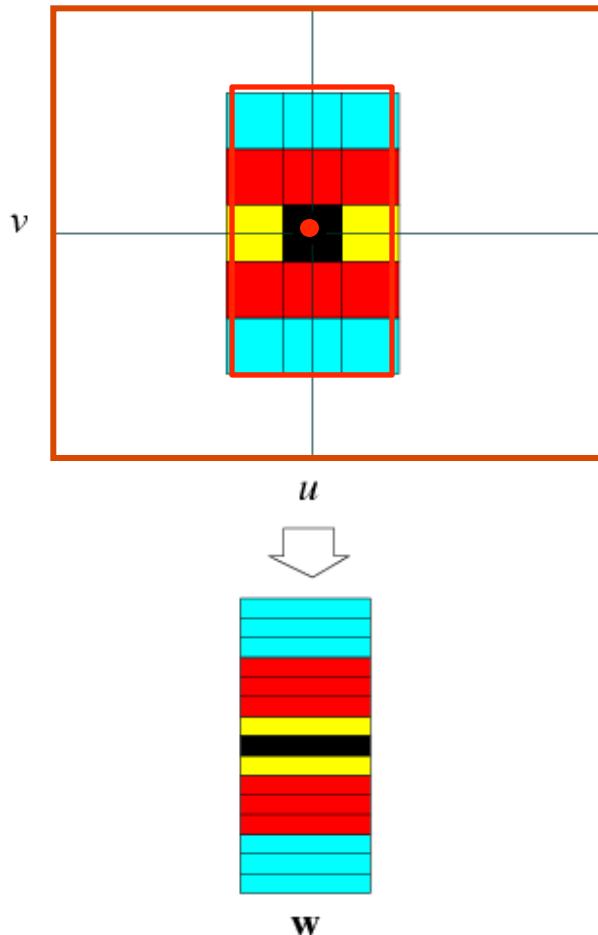


Correlation Methods (1970--)



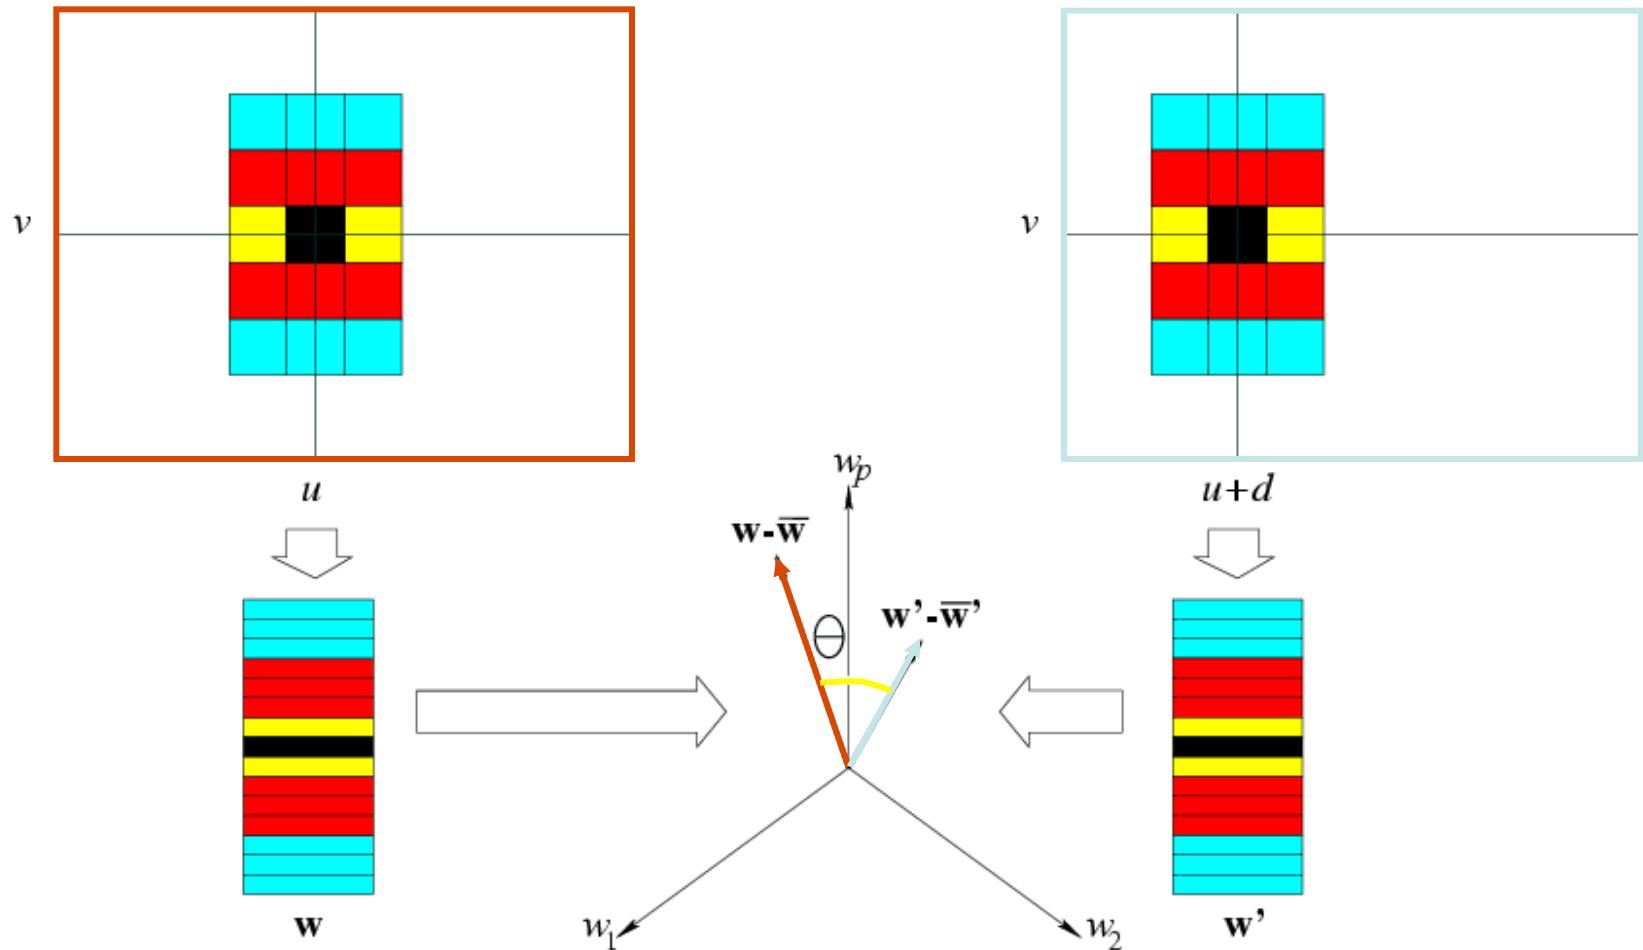
- Pick up a window around $p(u,v)$

Correlation Methods (1970--)



- Pick up a window around $p(u,v)$
- Build vector W
- Slide the window along v line in image 2 and compute w'
- Keep sliding until $w \cdot w'$ is maximized.

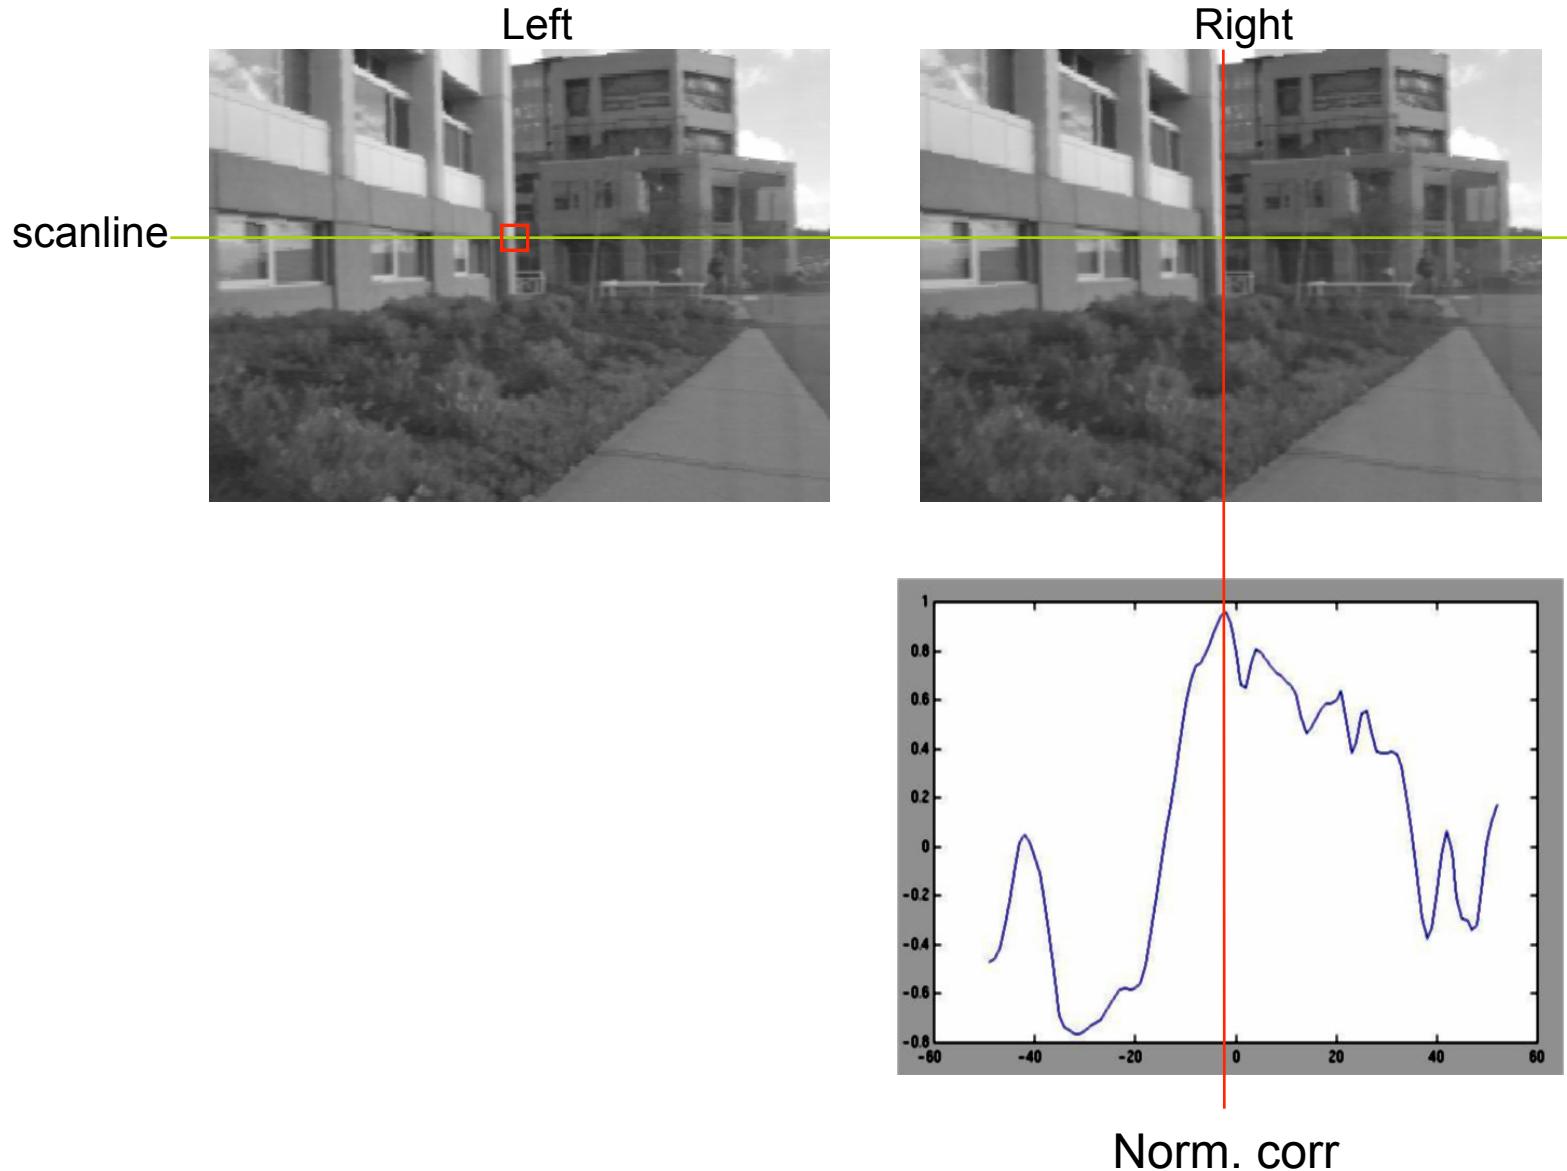
Correlation Methods (1970--)



Normalized Correlation; minimize:

$$\frac{(w - \bar{w})(w' - \bar{w}')}{\|(w - \bar{w})(w' - \bar{w}')\|}$$

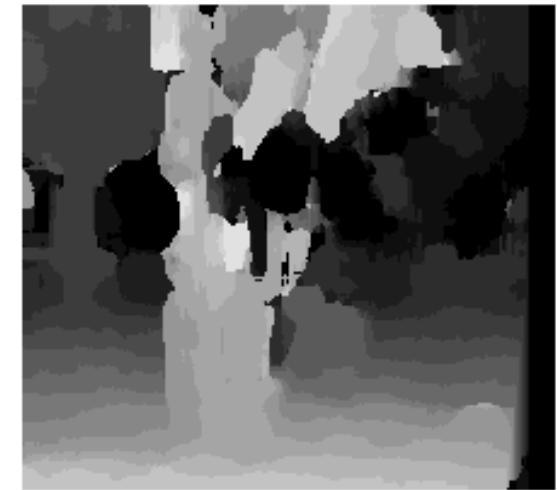
Correlation methods



Correlation methods



Window size = 3



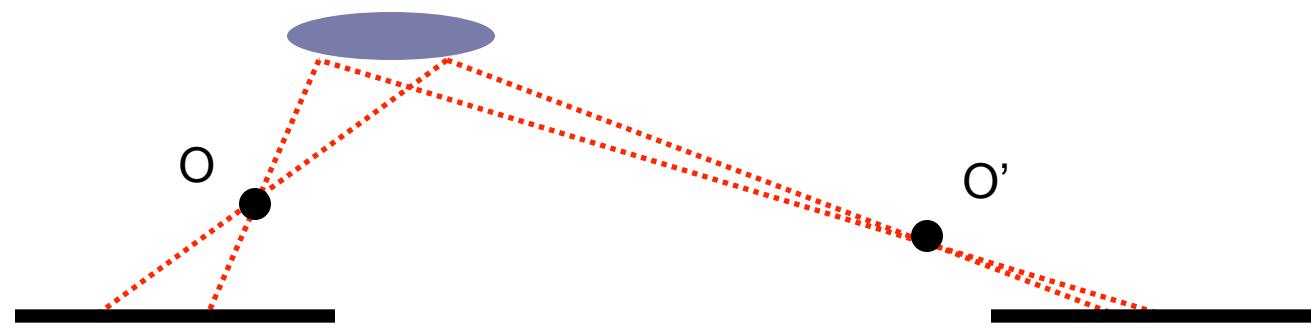
Window size = 20

- Smaller window
 - More detail
 - More noise

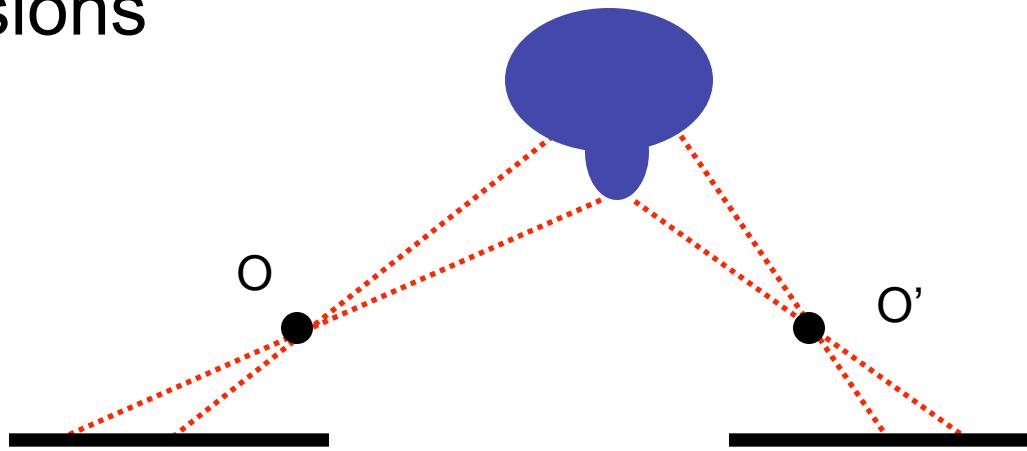
- Larger window
 - Smoother disparity maps
 - Less prone to noise

Issues

- Fore shortening effect

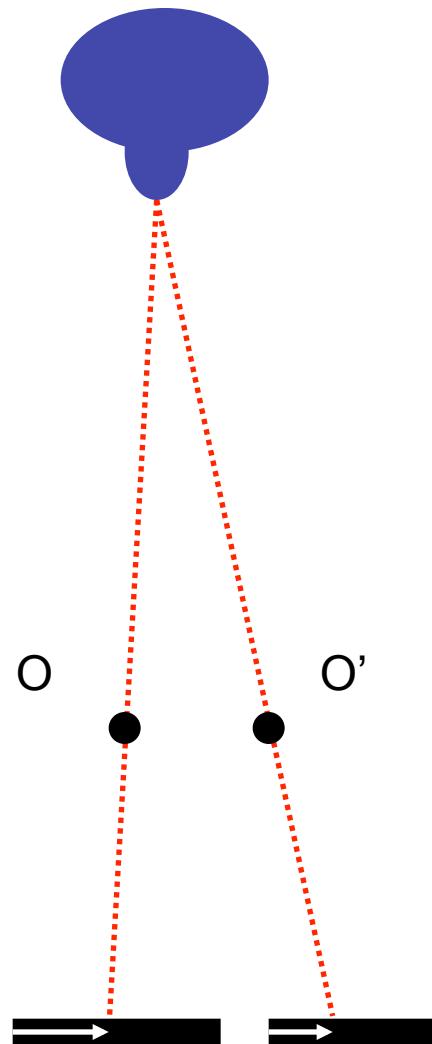


- Occlusions



Issues

- It is desirable to have small B/z ratio!
- Small error in measurements implies large error in estimating depth



Issues

- Homogeneous regions



Hard to match pixels in these regions

Issues

- Repetitive patterns



Correspondence problem is difficult!

- Occlusions
- Fore shortening
- Baseline trade-off
- Homogeneous regions
- Repetitive patterns

Apply non-local constraints to help enforce the correspondences

Results with window search

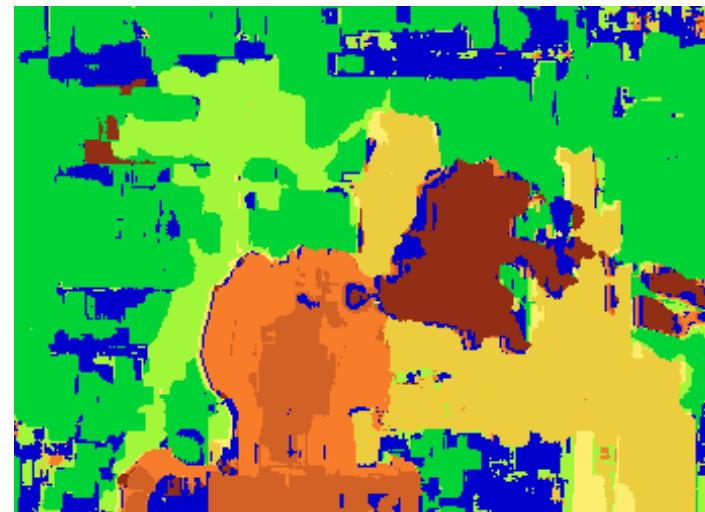
Data



Ground truth



Window-based matching

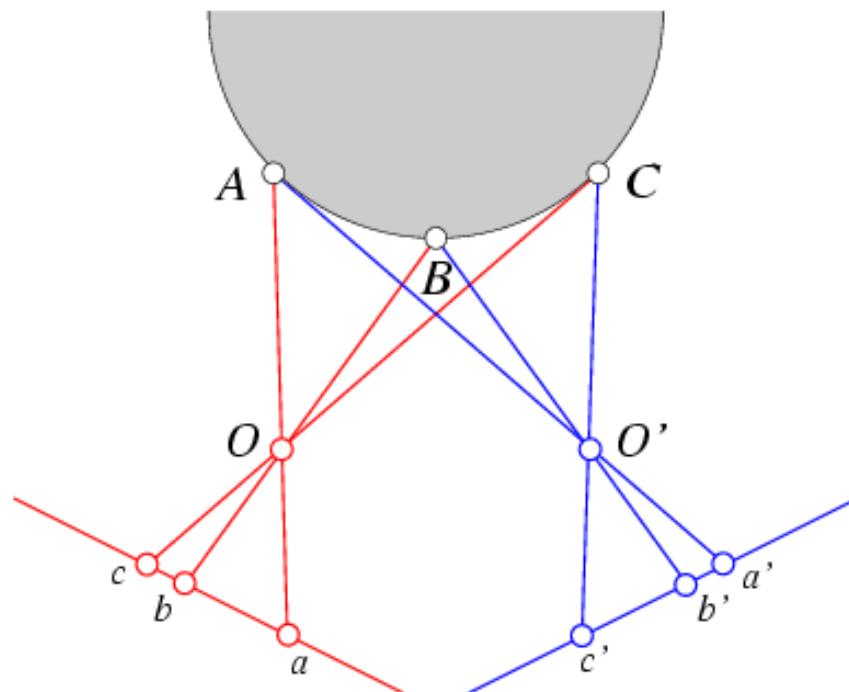


Improving correspondence: Non-local constraints

- Uniqueness
 - For any point in one image, there should be at most one matching point in the other image

Improving correspondence: Non-local constraints

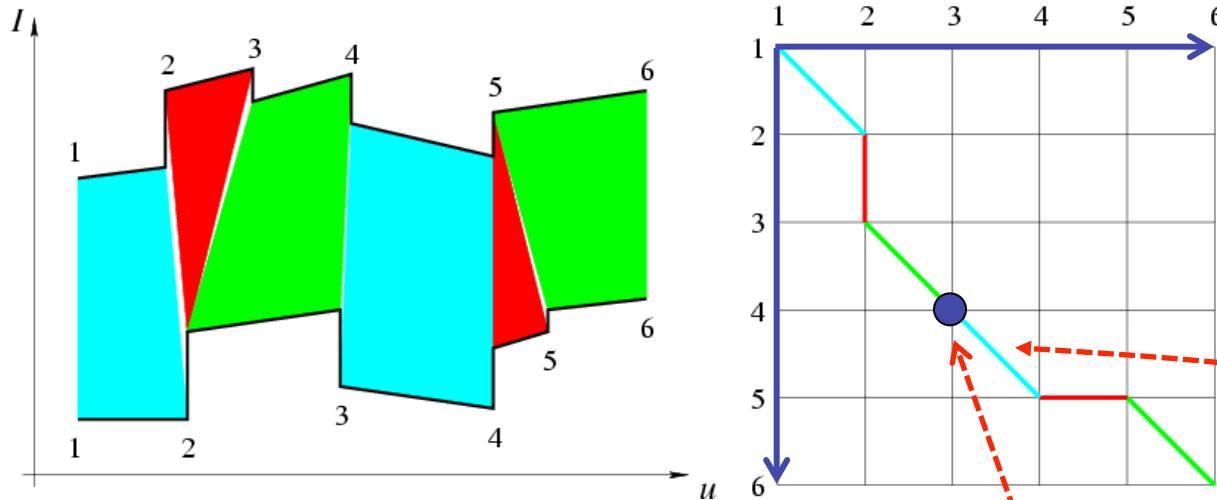
- Uniqueness
 - For any point in one image, there should be at most one matching point in the other image
- Ordering
 - Corresponding points should be in the same order in both views



Not always
in presence
of occlusions!

Dynamic Programming (Baker and Binford, 1981)

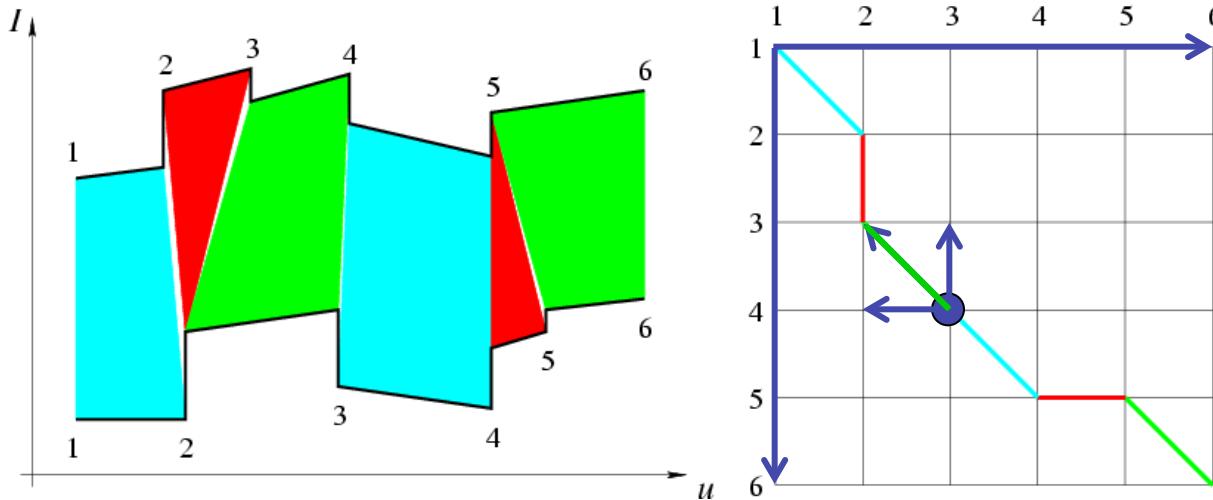
[Uses ordering constraint]



- Nodes = matched feature points (e.g., edge points).
- Arcs = matched intervals along the epipolar lines.
- Arc cost = discrepancy between intervals.

Find the minimum-cost path going monotonically down and right from the top-left corner of the graph to its bottom-right corner.

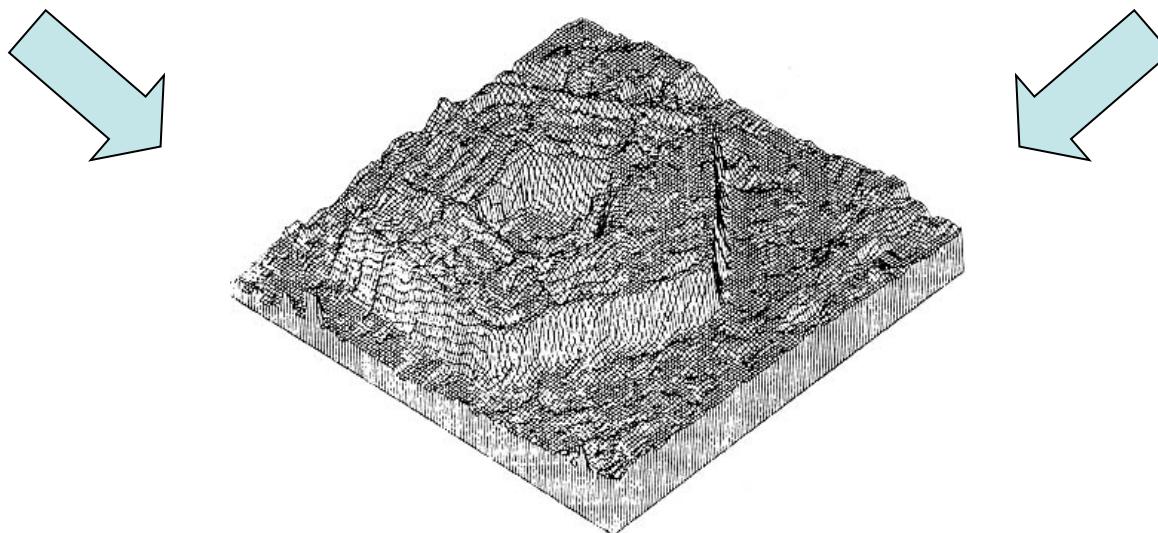
Dynamic Programming (Baker and Binford, 1981)



```
% Loop over all nodes  $(k, l)$  in ascending order.
for  $k = 1$  to  $m$  do
    for  $l = 1$  to  $n$  do
        % Initialize optimal cost  $C(k, l)$  and backward pointer  $B(k, l)$ .
         $C(k, l) \leftarrow +\infty$ ;  $B(k, l) \leftarrow \text{nil}$ ;
        % Loop over all inferior neighbors  $(i, j)$  of  $(k, l)$ .
        for  $(i, j) \in \text{Inferior - Neighbors}(k, l)$  do
            % Compute new path cost and update backward pointer if necessary.
             $d \leftarrow C(i, j) + \text{Arc - Cost}(i, j, k, l)$ ;
            if  $d < C(k, l)$  then  $C(k, l) \leftarrow d$ ;  $B(k, l) \leftarrow (i, j)$  endif;
            endfor;
        endfor;
    endfor;
% Construct optimal path by following backward pointers from  $(m, n)$ .
 $P \leftarrow \{(m, n)\}$ ;  $(i, j) \leftarrow (m, n)$ ;
while  $B(i, j) \neq \text{nil}$  do  $(i, j) \leftarrow B(i, j)$ ;  $P \leftarrow \{(i, j)\} \cup P$  endwhile.
```

courtesy slide to J. Ponce

Dynamic Programming (Ohta and Kanade, 1985)

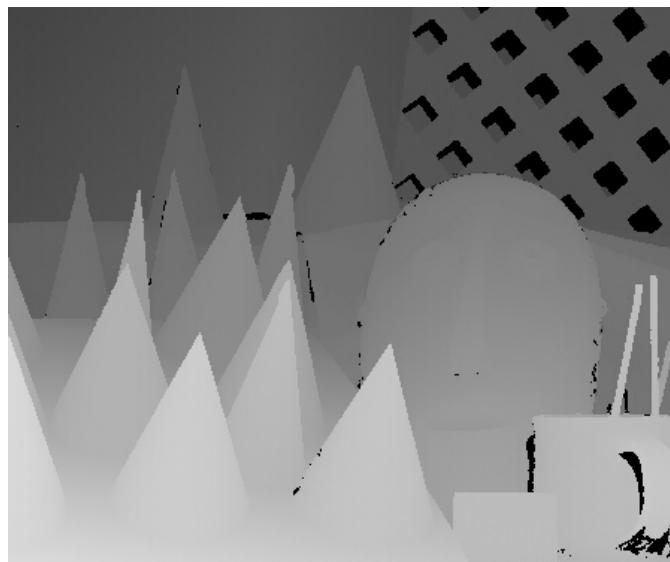
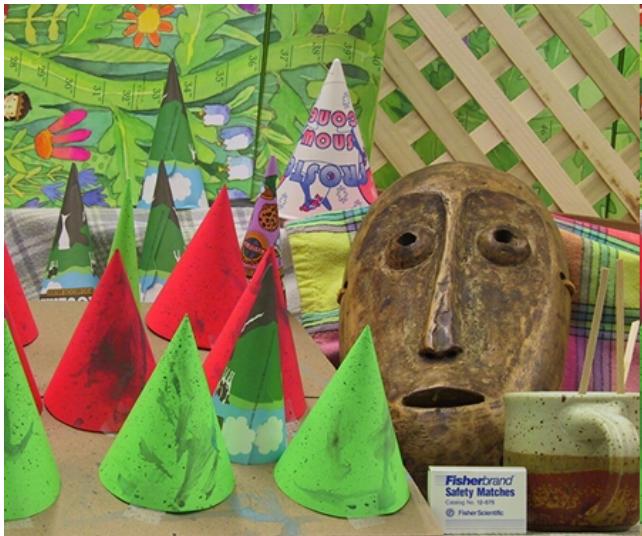


Reprinted from "Stereo by Intra- and Intet-Scanline Search," by Y. Ohta and T. Kanade, IEEE Trans. on Pattern Analysis and Machine Intelligence, 7(2):139-154 (1985). © 1985 IEEE.

Improving correspondence: Non-local constraints

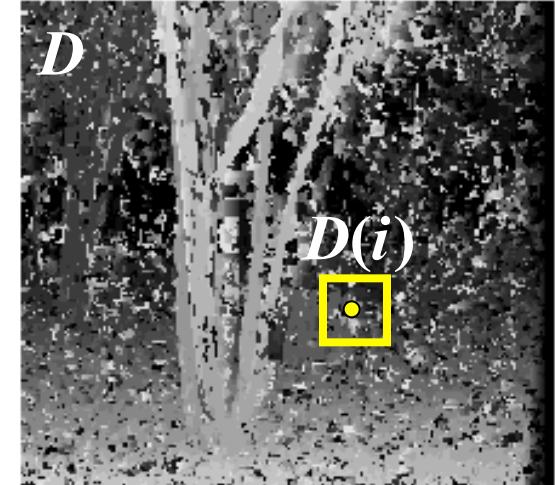
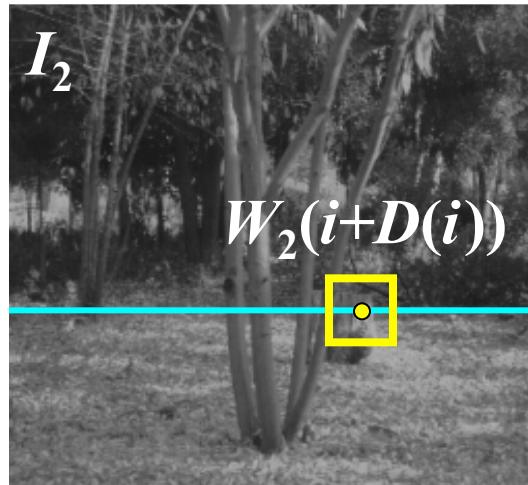
- Uniqueness
 - For any point in one image, there should be at most one matching point in the other image
- Ordering
 - Corresponding points should be in the same order in both views
- Smoothness
 - Disparity is typically a smooth function of x (expect in occluding boundaries)

Smoothness



Stereo matching as energy minimization

Y. Boykov, O. Veksler, and R. Zabih, Fast Approximate Energy Minimization via Graph Cuts, PAMI 01



$$E = \alpha E_{\text{data}}(I_1, I_2, D) + \beta E_{\text{smooth}}(D)$$

$$E_{\text{data}} = \sum_i (W_1(i) - W_2(i + D(i)))^2$$

$$E_{\text{smooth}} = \sum_{\text{neighbors } i, j} \rho(D(i) - D(j))$$

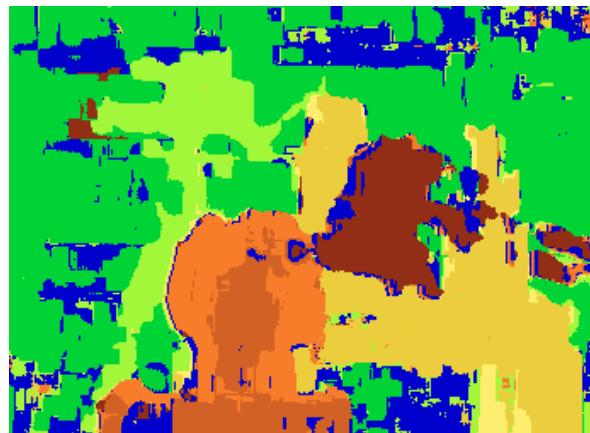
- Energy functions of this form can be minimized using *graph cuts*

Stereo matching as energy minimization

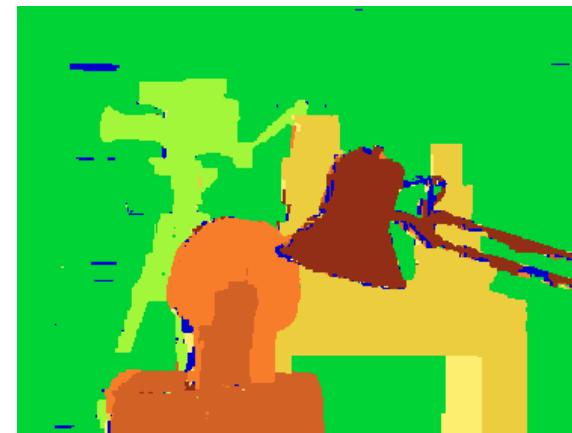
Y. Boykov, O. Veksler, and R. Zabih, Fast Approximate Energy Minimization via Graph Cuts, PAMI 01



Ground truth



Window-based
matching



Graph cuts

Two-frame stereo correspondence algorithms

[Click here](#)

<http://www.middlebury.edu/stereo/>

Stereo SDK stereo vision software development kit.

A. Criminisi, A. Blake and D. Robertson

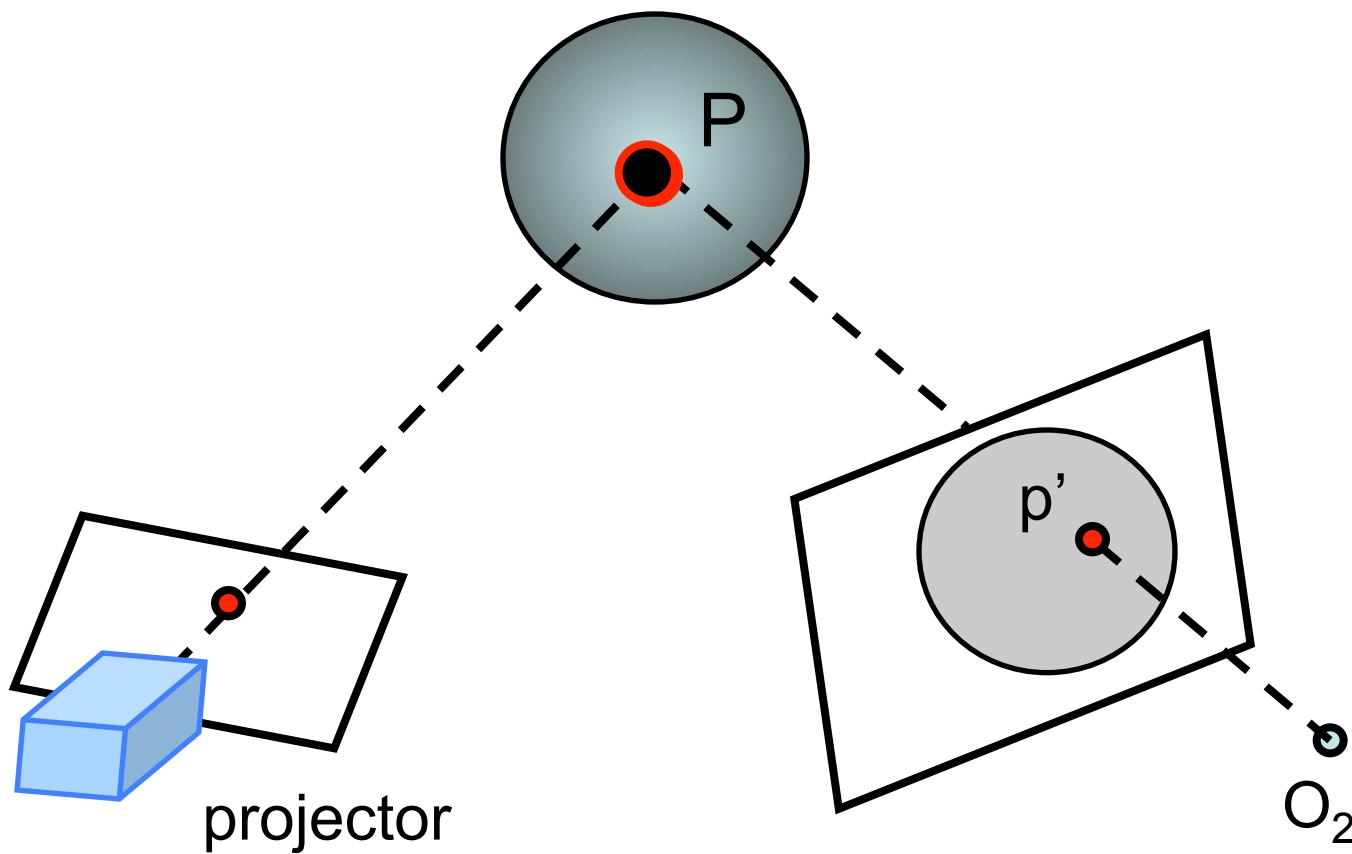




Stereo systems

- Stereo vision
- Rectification
- Correspondence problem
- Active stereo vision systems

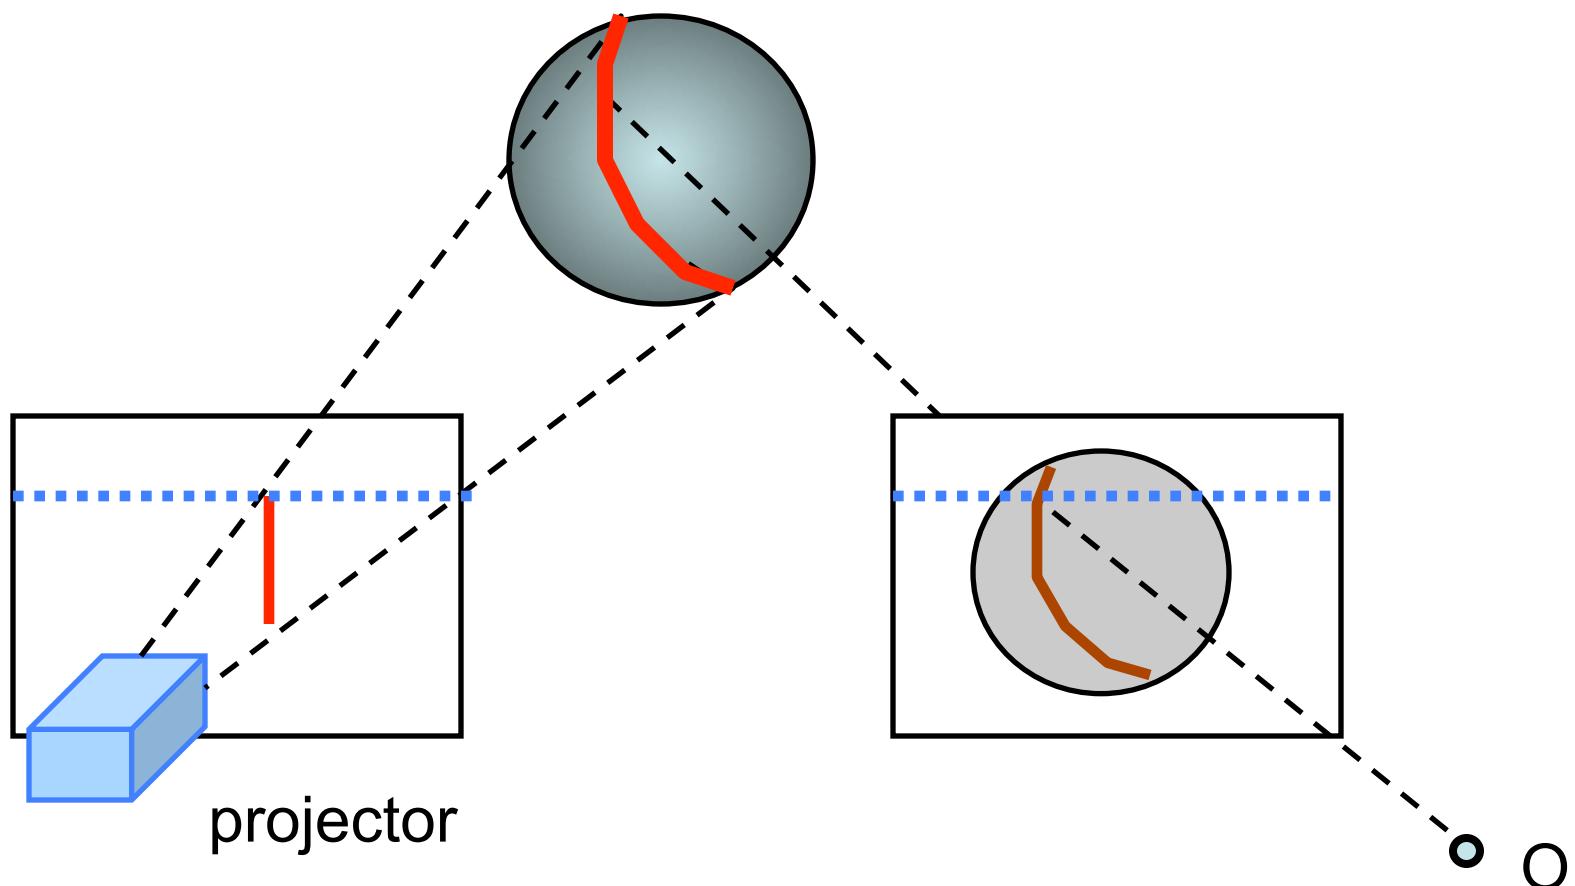
Active stereo (point)



Replace one of the two cameras by a projector

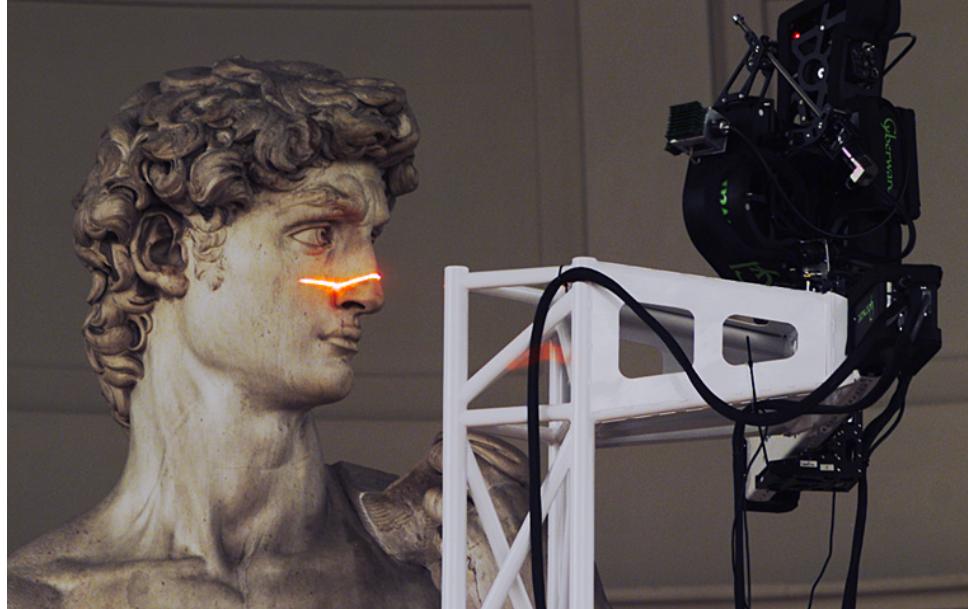
- Single camera
- Projector geometry calibrated
- What's the advantage of having the projector? Correspondence problem solved!

Active stereo (stripe)



- Projector and camera are parallel
- Correspondence problem solved!

Laser scanning



Digital Michelangelo Project
<http://graphics.stanford.edu/projects/mich/>

- Optical triangulation
 - Project a single stripe of laser light
 - Scan it across the surface of the object
 - This is a very precise version of structured light scanning

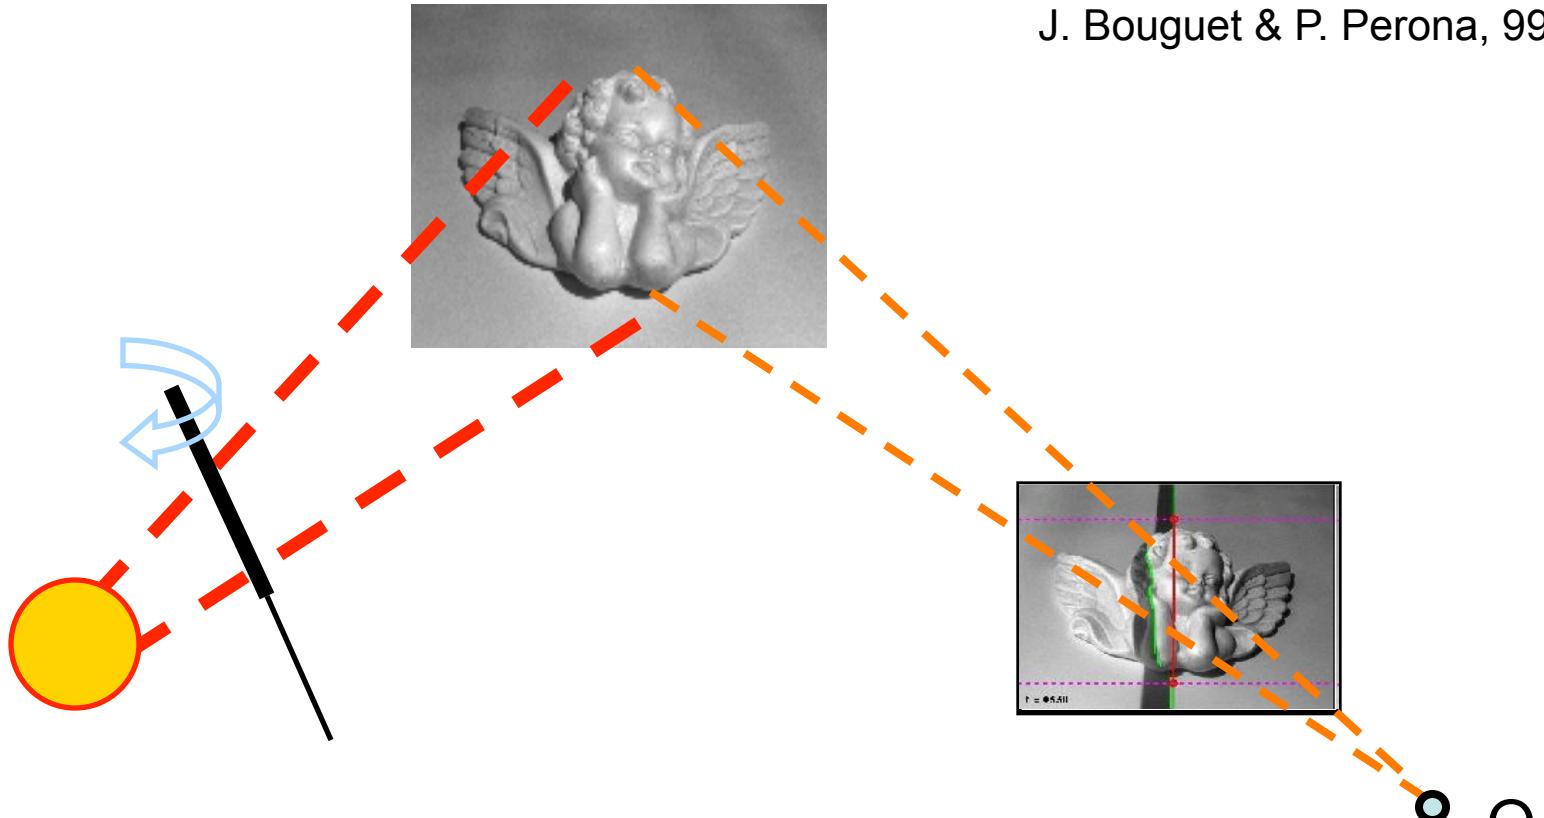
Laser scanning



The Digital Michelangelo Project, Levoy et al.

Source: S. Seitz

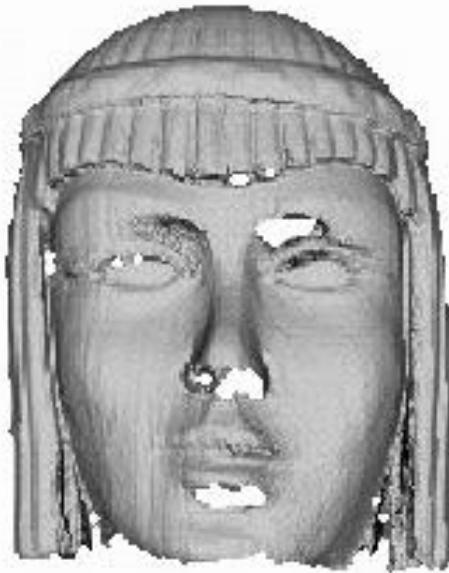
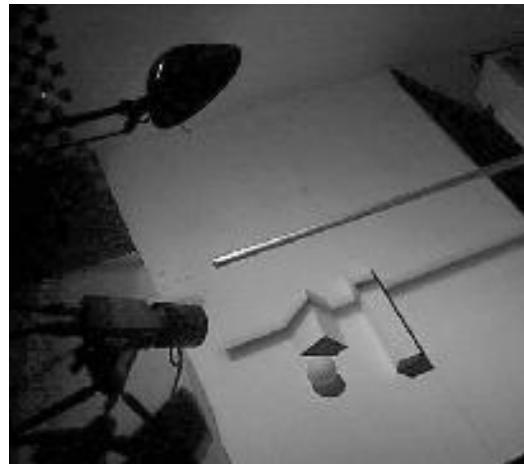
Active stereo (shadows)



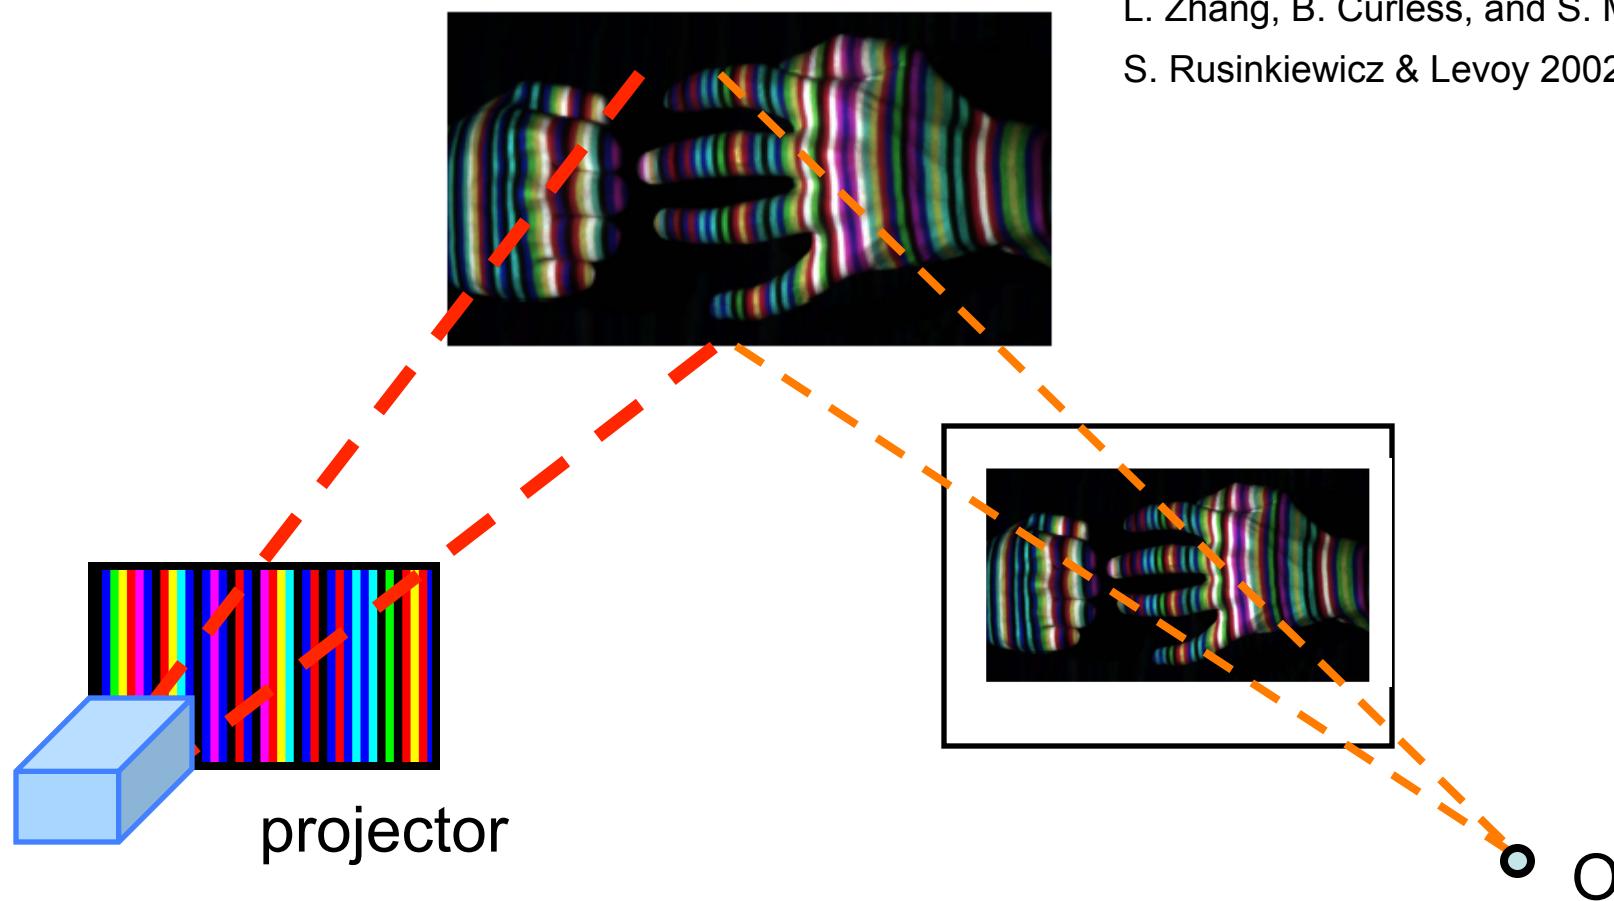
Light source

- 1 camera, 1 light source
- very cheap setup
- calibrated the light source

Active stereo (shadows)

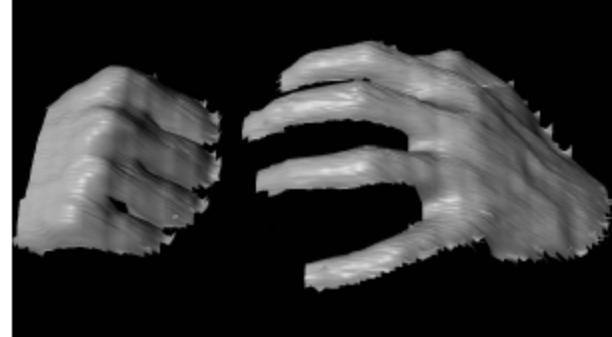
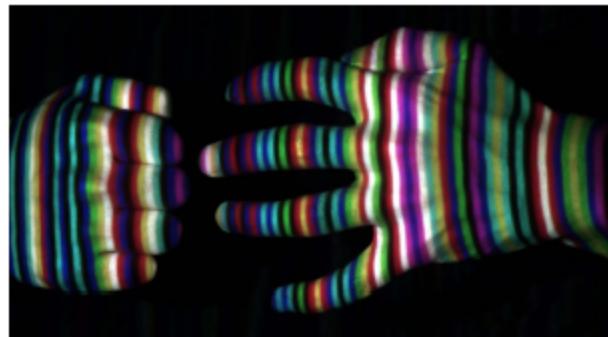


Active stereo (color-coded stripes)

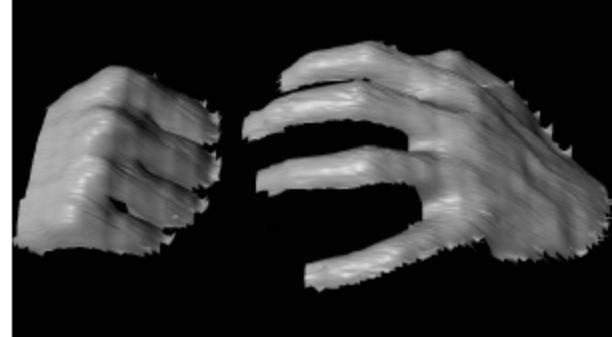


L. Zhang, B. Curless, and S. M. Seitz 2002
S. Rusinkiewicz & Levoy 2002

- Dense reconstruction
- Correspondence problem again
- Get around it by using color codes



L. Zhang, B. Curless, and S. M. Seitz. Rapid Shape Acquisition Using Color Structured Light and Multi-pass Dynamic Programming. *3DPVT* 2002



L. Zhang, B. Curless, and S. M. Seitz. Rapid Shape Acquisition Using Color Structured Light and Multi-pass Dynamic Programming. *3DPVT* 2002

Rapid shape acquisition: Projector + stereo cameras



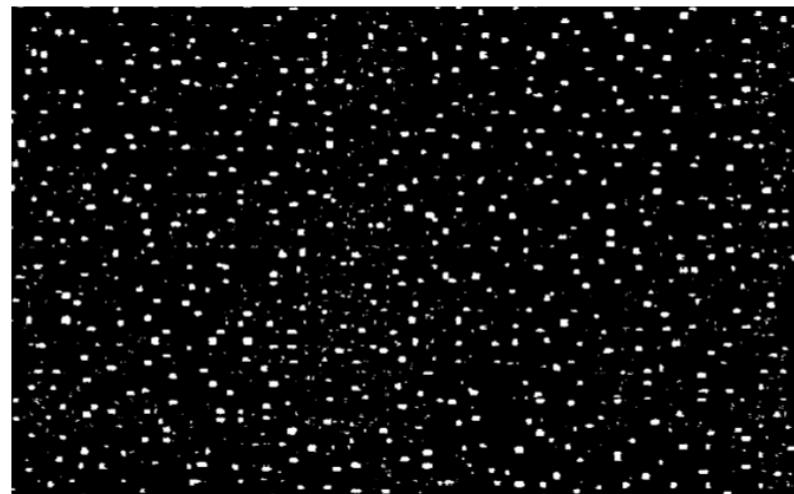
Microsoft Kinect

The Kinect combines structured light with two classic computer vision techniques: depth from focus, and depth from stereo.

Stage 1: The depth map is constructed by analyzing a speckle pattern of infrared laser light



The Kinect uses infrared laser light, with a speckle pattern



Shpunt et al, PrimeSense patent application
US 2008/0106746

Next lecture...

Affine Structure from Motion

Human Stereopsis

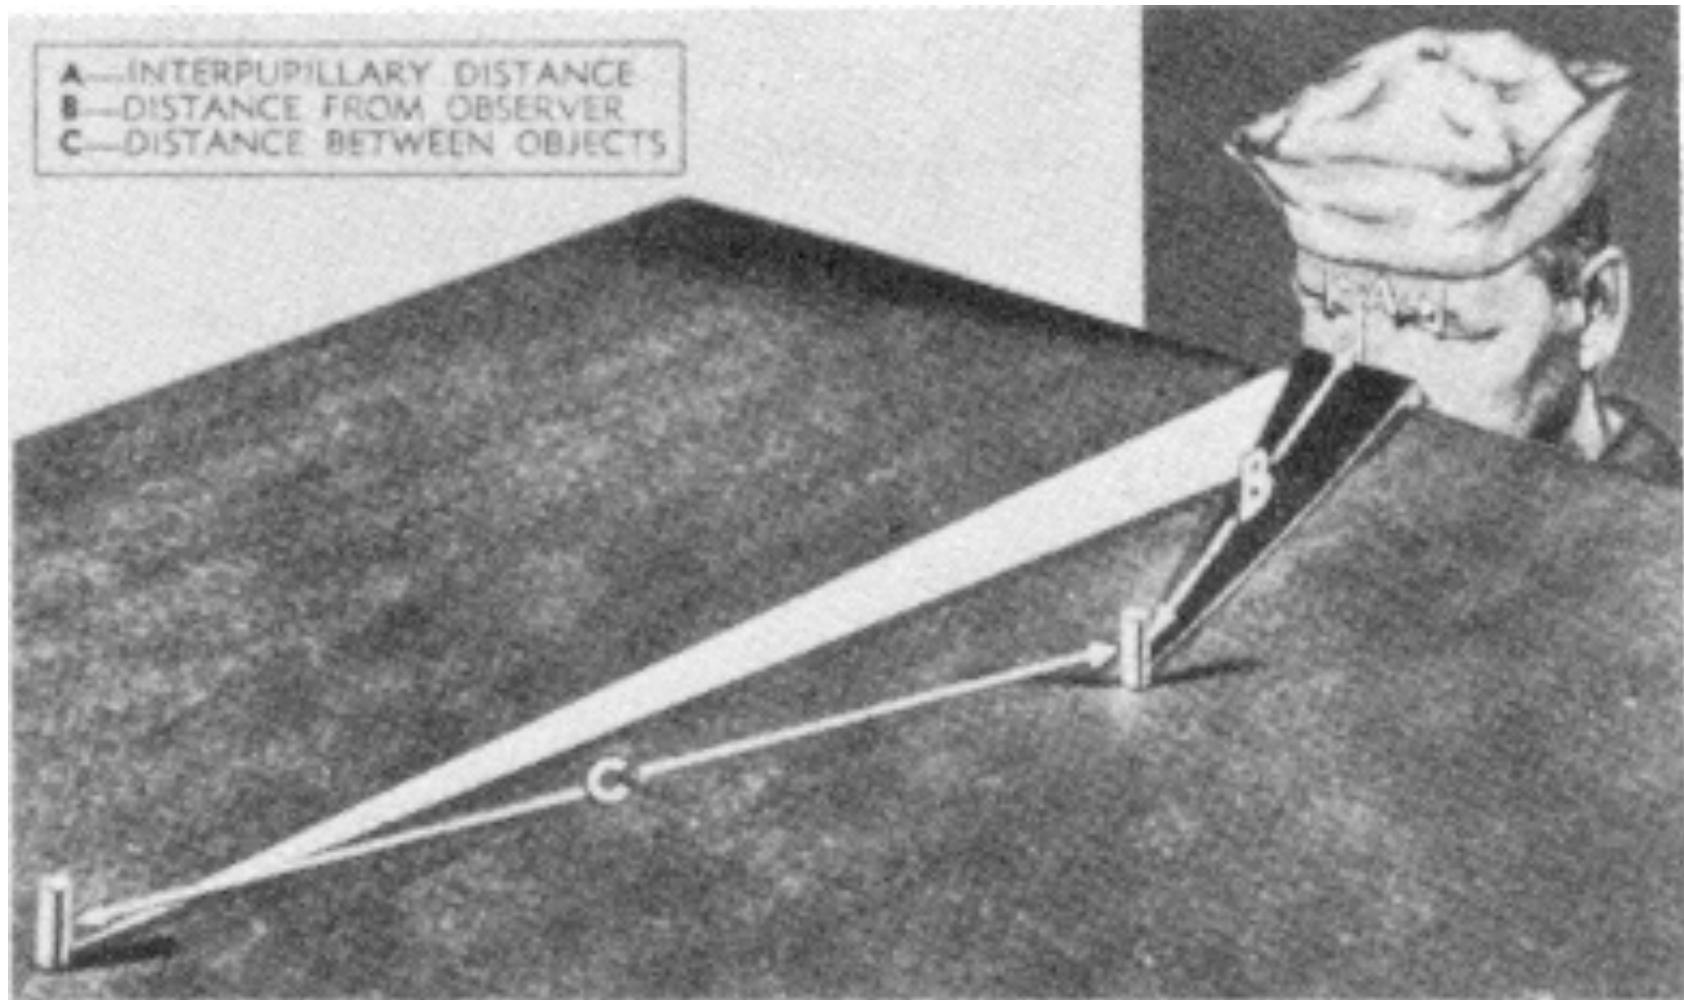
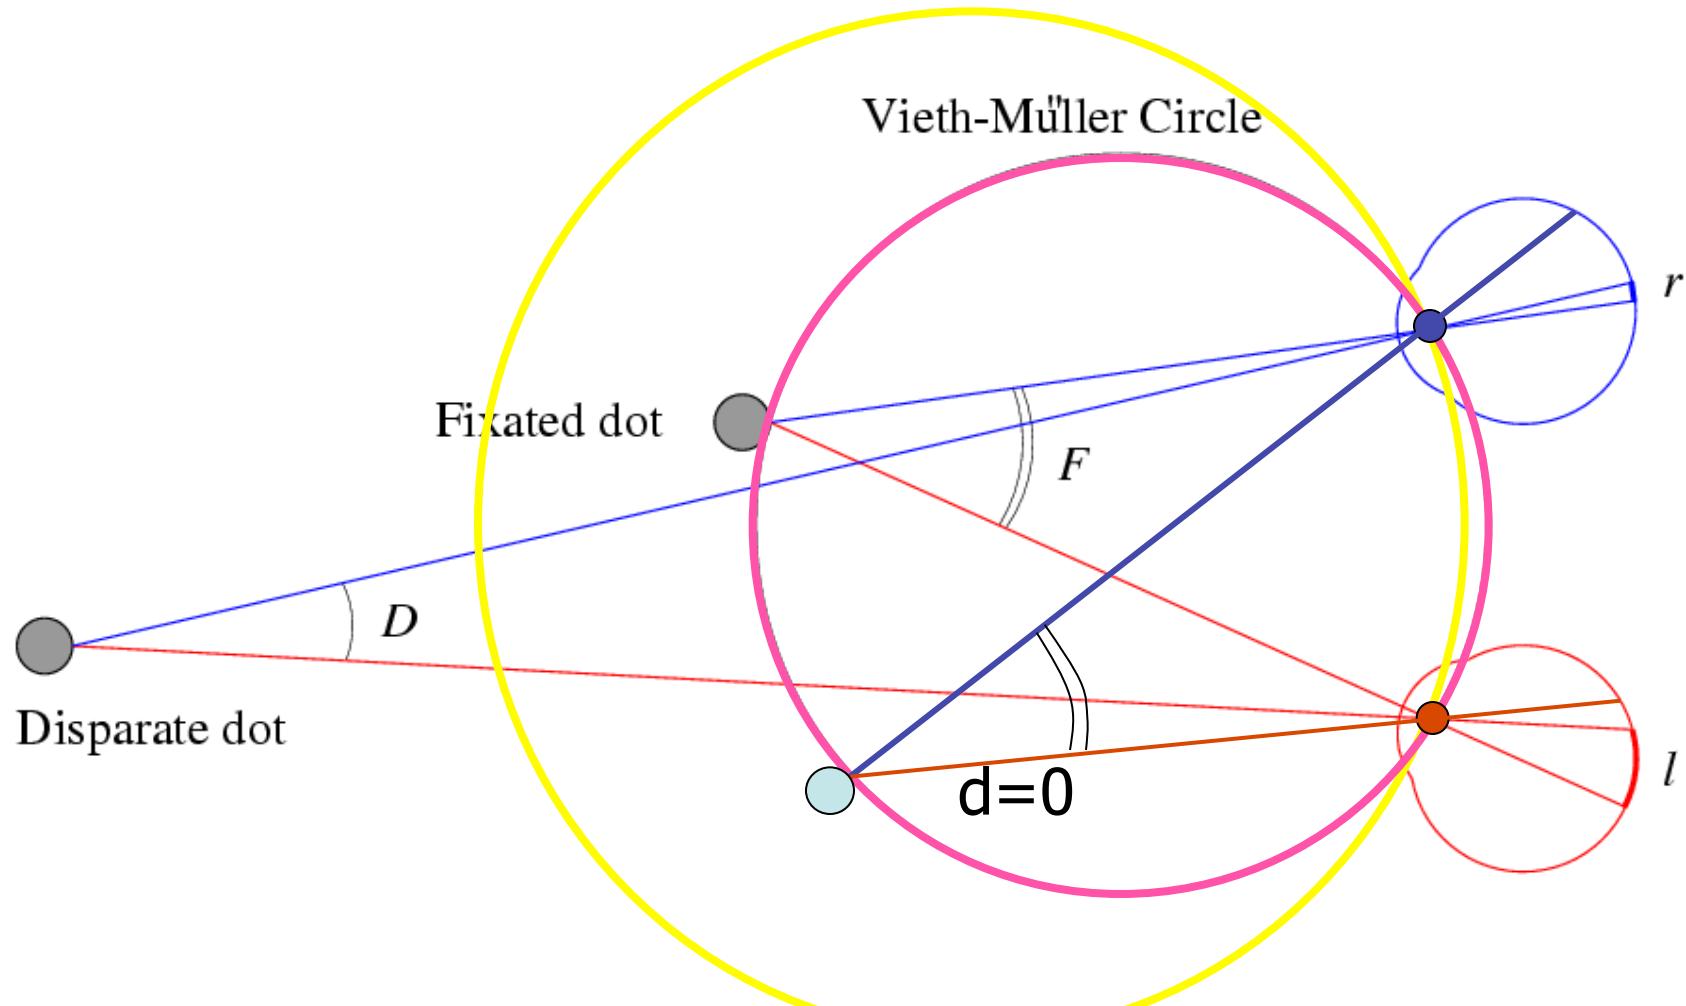


Figure from US Navy Manual of Basic Optics and Optical Instruments, prepared by Bureau of Naval Personnel. Reprinted by Dover Publications, Inc., 1969.

Human Stereopsis: Reconstruction



$$\text{Disparity: } d = r-l = D-F; \quad d < 0$$

In 3D, the horopter.

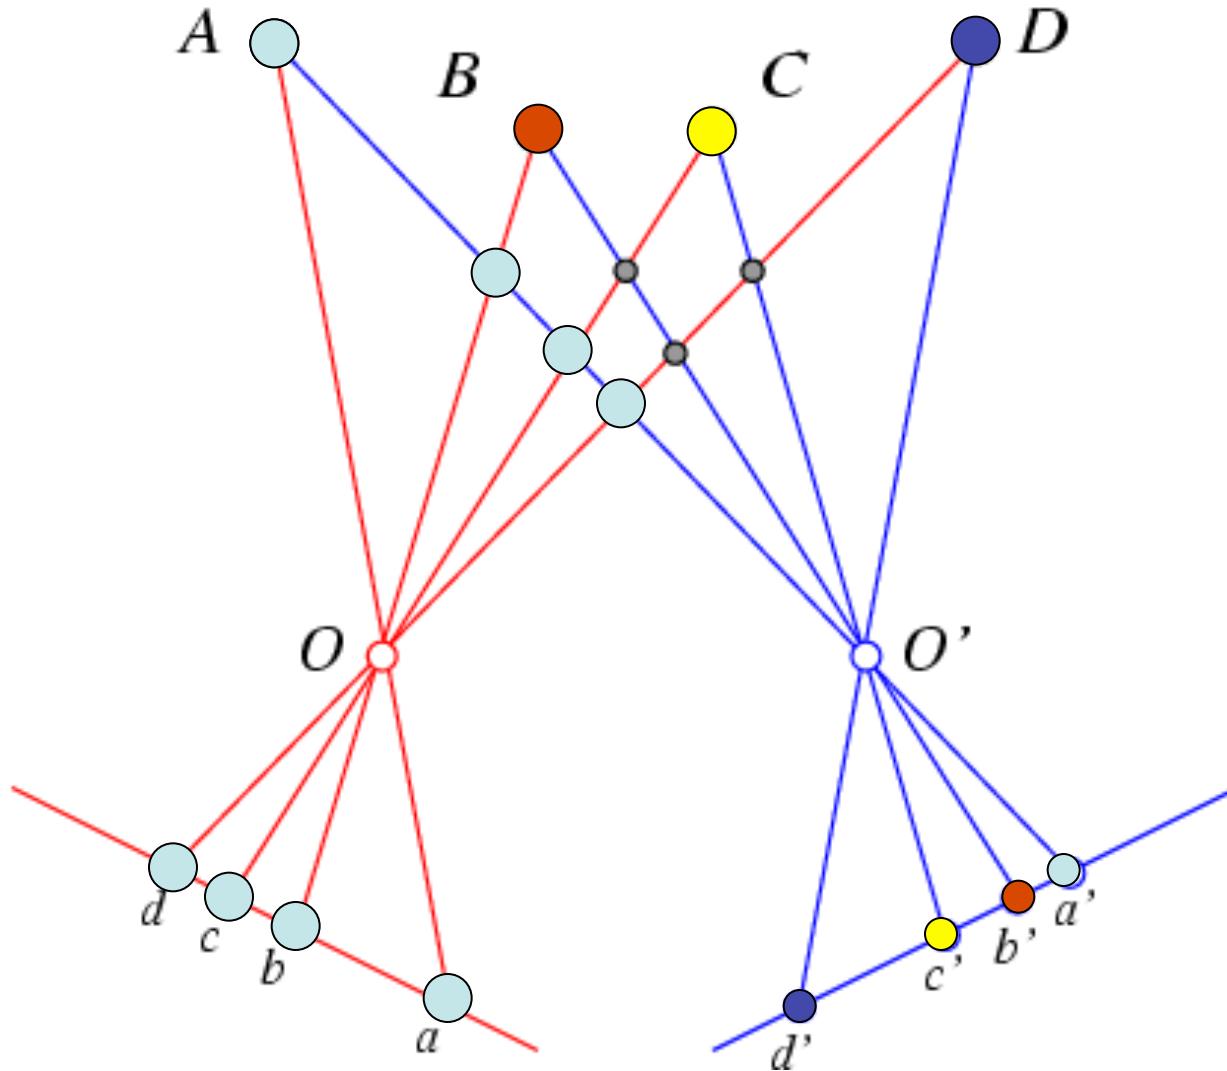
Human Stereopsis: Reconstruction

What if F is not known?

Helmoltz (1909):

- There is evidence showing the vergence angles cannot be measured precisely.
- Humans get fooled by bas-relief sculptures.

Human Stereopsis: Binocular Fusion

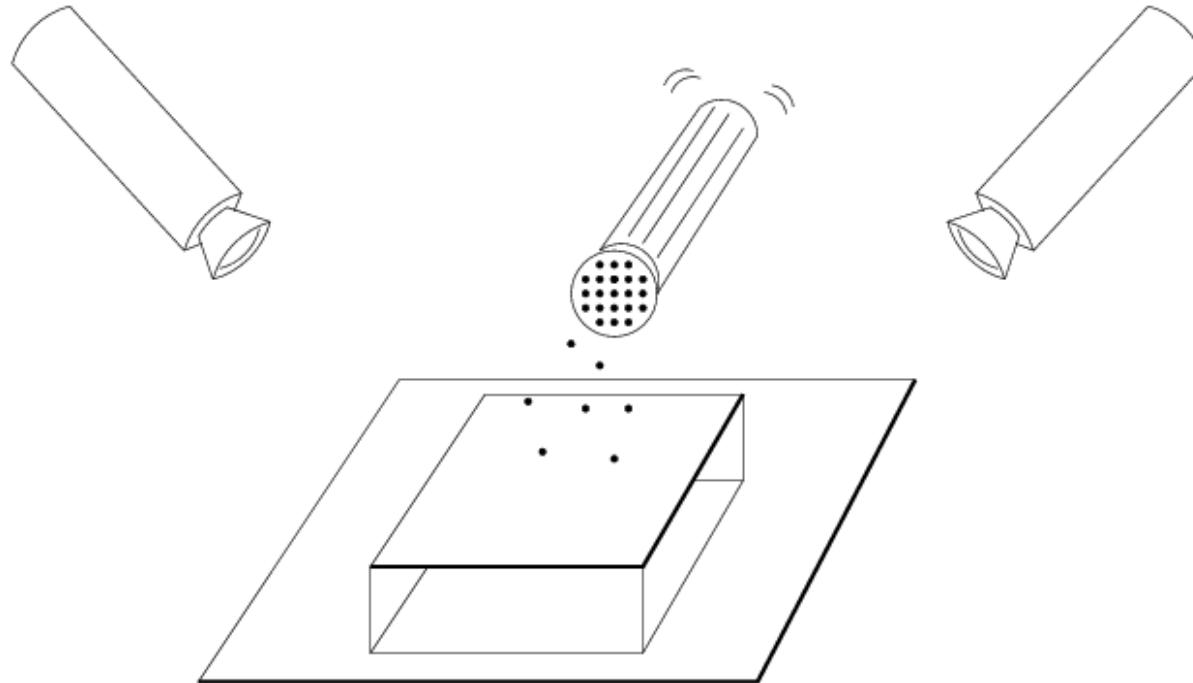


Human Stereopsis: Binocular Fusion

How are the correspondences established?

Julesz (1971): Is the mechanism for binocular fusion a monocular process or a binocular one??

- There is anecdotal evidence for the latter (camouflage).



- Random dot stereograms provide an objective answer

Issues

